

Scraper

A daily industry news desk that reads 35 outlets so you read twelve stories.

Scraper is the ecosystem's **fresh third-party blood**. Every weekday morning it sweeps the trade press and the federal record, scores every article against what *you* actually cover, clusters duplicate reporting into single stories, writes a summary and a separately-labeled analytical read, and delivers two editions before you sit down. Everything it discards is **counted and explained** — the footer never claims nothing was held back.

35	29	2	6 mo
CURATED OUTLETS	TUNABLE SEGMENTS	DAILY EDITIONS	SEARCHABLE ARCHIVE

INSIDE

- 01

What Scraper does
- 02

The source roster & its discipline

How an article is scored

The morning build pipeline
- 03

Your day with Scraper

Tuning & the feedback loop

Getting the most out of it
- 04

The corpus bridge

Two editions, built before you wake up

Scraper's job is not to show you the news. It is to decide, defensibly, which twelve of the day's hundred-odd articles are worth your attention — and to show its working when you disagree.

The morning digest

Two editions build daily at 6am desk time and deliver from 7am: **Your Morning Digest (BESS)**, tuned to utility-scale storage, and **Your Morning Digest (AIDC)**, tuned to the data-center power chain. Each edition sorts what survives into three sections:

COMPANIES	MARKET	INCIDENTS
Stories naming a company you cover in Profiler — the ones that change a dossier or a conversation.	Structural movement with no single company at the centre: pricing, policy, interconnection, demand.	Fires, safety findings, siting opposition and permitting fights — the stories that move public sentiment.

Every item carries a reported **summary** and, set apart from it, **the desk's read** — an analytical line kept deliberately shorter than the reporting. The two are colour-keyed and the key is printed in the footer: **amber means analysis**, and concrete figures are highlighted so a number never hides inside a paragraph.

The rest of the app

SURFACE	WHAT IT GIVES YOU
Tune	29 segments and 24 topics with individual on/off toggles and weights, plus your covered companies synced live from Profiler. This is where you tell the scorer what you care about today.
Editions	Create your own beyond the two defaults — name it, set the cadence and window, apply a tuning preset, schedule it.
Archive	About six months of past editions, browsable as a grid, a table, or a calendar. Links keep resolving to their real destinations for roughly sixteen weeks.
Projects & Plan	Saved research queries with backfill, deep backfill and compile — for going wide on a topic rather than waiting for it to surface.
Subscribers	An email list, "email me the latest", and shareable links for a single edition.
Why thin? / Held back	Two explanations rather than a shrug: why an edition is short today, and exactly what cleared the bar but did not make the cut.
Source stats	Per-outlet contribution — who is actually earning their place in the roster.
Corpus route	A read-only, token-gated feed that hands scored articles and relationship candidates to Profiler and Classroom.

The roster, and why it is short on purpose

Thirty-five outlets is a deliberate number. A larger roster does not produce a better digest — it produces more noise for the scorer to reject, and more silent failures nobody notices.

Tier 1 — the record and the beat leaders

Seventeen feeds, including five **primary federal sources** adopted after an executive order naming battery storage and grid inverters as in-scope critical infrastructure reached the desk only through second-hand trade coverage, days late and partial. The roster had been entirely secondary — thirty outlets reporting on government action, none reading it.

PRIMARY SOURCE	WHY IT IS THE RECORD, NOT A REPORT OF IT
White House — Presidential Actions	Executive orders in full. The operative detail is in the definitions section that trade pieces summarise away.
Federal Register — FERC	A FERC order takes legal effect on publication here, not on the agency's own newsroom.
Federal Register — IRS	Credit and domestic-content guidance — the rules that move project economics.
US Dept. of Energy	Programme announcements, loan authority, funding awards.
EIA — Today in Energy	The primary data series behind most market claims you will read elsewhere.

The rest of Tier 1 is the beat leadership — Utility Dive, Data Center Dynamics, Energy-Storage.news, ESS News, Canary Media, POWER, Data Center Knowledge, Data Center Frontier, pv magazine USA, Microgrid Knowledge, T&D World, The Next Platform. Tier 2 adds eighteen outlets of breadth and specialist depth.

NO OUTLET IS ADOPTED FROM REPUTATION

Every feed is **probed live** before it joins the roster: HTTP status, that the body is actually XML, and that it carries items with recent dates. A feed URL is never taken from memory or from the outlet's standing in the industry. When one is dropped, the reason and the date are recorded so a future session cannot re-propose it as "a gap in coverage" — which is precisely what used to happen.

Blocked is not the same as gone

Two failure modes look identical in a browser and are opposite facts. Scraper distinguishes them and treats them differently in the UI:

Blocked — the outlet is alive

A bot wall answers every automated client. Kept visible in Tune, struck through, toggle frozen off — so you can see the beat is *intentionally* uncovered rather than quietly missing.

Offline — the publication ended

The domain is parked and the feed is a redirect stub. Filtered out entirely; the record is kept so re-adding the key would reactivate it if it ever returns.

Six outlets currently sit on the do-not-re-propose list, five blocked and one genuinely gone. None of them cost coverage: battery and storage run through Energy-Storage.news and ESS News, solar through pv magazine USA and Solar Power World, data centres through Data Center Frontier, DCD and The Register, and the federal agencies through the Federal Register — where their output is legally effective anyway.

How an article is scored

Every article gets a number out of 100 from four weighted bands, then a geographic multiplier. Below 25 it does not even enter the intake — it is stored for the corpus and never shown.

BAND	WEIGHT	WHAT EARNS IT
Company	40	Names a company you cover in Profiler. Matched on the dossier's permanent slug, not its display name, so a rename never silently detaches a row from its company. Additional companies add a smaller increment each.
Topic	25	Hits one of 24 topic vocabularies — interconnection, tariffs, offtake, safety, siting, supply chain and the rest.
Substance	20	Carries concrete detail — figures, capacities, dates, named counterparties — rather than a press-release restatement.
Emphasis	15	Weighted toward companies you are actively working: a dossier updated within the last 45 days earns the full recency component.

The segment gate — specificity beats breadth

Segments are hierarchical. `seg-bess` is the broad parent; beneath it sit utility-scale, data-centre/behind-the-meter, residential, C&I and long-duration. Two rules follow from that shape:

- **A child hit demotes its parent.** When an article matches a specific child segment, the generic parent's evidence is discounted — a residential product story cannot ride the broad "energy storage" match into a utility-scale edition.
- **Off-segment company news is zeroed, not merely down-ranked.** If an article's only segment hits are segments you have turned *off*, its company, topic and engagement signals are all set to zero. A covered automaker's vehicle recall no longer clears the bar on the strength of the company match when you have "EVs & automotive" switched off.
- **No segment hit at all changes nothing.** The gate only fires on positive evidence of an excluded subject — never on absence.

Geography — a multiplier, not a penalty band

The US market is priority one. Related markets — China, Mexico, Chile, Canada — matter mainly when what happened there affects the US. That is applied as a multiplier on the finished score, because an additive penalty would leave a strong company match on a foreign story still clearing the bar:

×1.00 — US OR UNMARKED	×0.55 — RELATED MARKET	×0.25 — ELSEWHERE
The load-bearing default. Most US trade coverage never says "United States" — it says ERCOT, or a county in Texas, or nothing at all.	A foreign subject in a priority-two country with no US connection named anywhere in the piece.	Everywhere else. The lead should never be an Australian transmission story.

WHY THE DEFAULT IS THE IMPORTANT PART

Penalising articles with *no* geographic marker would empty the digest, because the overwhelming majority of relevant US coverage carries no such marker. A penalty therefore requires positive evidence of a foreign subject — never merely the absence of evidence of a domestic one. This is the single rule most likely to be broken by a well-meaning "improvement" to the scorer.

The morning build

The build runs in chunked steps rather than one long pass, so a slow feed or a rate-limited model degrades one step instead of failing the edition. Each phase hands its state to the next.

- 1 Fetch**
Walk the enabled feeds for the edition's window, parse each one, and drop anything already seen. A feed that fails is skipped and logged — never allowed to abort the run.
- 2 Score & gate**
Apply the rubric, the segment gate and the geographic multiplier to every item. Anything below 25 is stored for the corpus and goes no further.
- 3 Cluster**
Detect articles covering the same event and collapse them into one story, keeping the best-ranked outlet's version. Six outlets reporting one deal is one item, not six.
- 4 Summarize**
Generate the reported summary and, separately, the desk's analytical read — capped shorter than the reporting on purpose. Gemini or Claude Sonnet, selectable in the UI; failures are classified as soft (retry) or terminal (stop asking) so one bad call cannot burn the whole budget.
- 5 Backstop**
If the edition is coming in thin, a second pass widens the net rather than shipping three items with no explanation.
- 6 Render & deliver**
Build the HTML, resolve every link through the intake table so clicks land on the real article, store the edition, and email subscribers from 7am.
- 7 Reconcile**
A daily pass mines articles naming two covered companies into relationship candidates for Profiler, and prunes orphaned intake rows.

The footer tells the truth

Every edition ends with a line like `14 of 15 relevant · 104 scanned · 1 more held back`. That is not decoration — it is the contract. Items that cleared the relevance bar but lost to a section cap are counted in that number and readable in the **Held back** panel.

WHAT HAPPENS WHEN AN EDITION IS TOO BIG TO STORE

A Sheets cell caps at 50,000 characters, so a very large edition has to be trimmed to fit. The trim is ordered: analysis text goes first, then summaries, then the held-back list is halved. Crucially, when the held-back list is trimmed the edition **says so** rather than reporting a smaller number — because "nothing was held back" would not be a smaller truth, it would be a false one. The total count survives the trim untouched.

Retention

400 editions kept

Roughly six months of three weekday editions, browsable and searchable. Narrow reads mean depth no longer costs page-load time.

240 editions of live links

About sixteen weeks. Beyond that an edition stays readable but its links fall back to opening the app.

Your day with Scraper

The daily loop takes five minutes. The tuning loop runs weekly at most. Both feed the same scorer, and the second is what makes the first keep getting better.

The five-minute morning

- 1 **Open the edition**
It was built at 6am and is waiting. Read by section — Companies first if you are prepping a call, Market first if you are thinking about positioning.
- 2 **Read the summary, then the read**
The reported facts and the desk's interpretation are visually separated so you always know which is which. Figures are highlighted; you can scan for numbers alone.
- 3 **Click through what matters**
This is not passive. Clicks are logged and become an engagement boost for that company and topic, so tomorrow's ranking reflects today's attention.
- 4 **Rate the misses**
👍 and 👎 on individual items. A thumbs-down on a story that should not have been there is the highest-value signal you can give.
- 5 **Ask "Why thin?" when it looks light**
You get the real reason — a quiet news day, a failing feed, a segment you turned off — instead of guessing.
- 6 **Check "Held back"**
Everything that cleared the bar but lost to a section cap. Worth a glance on days when the top of the edition felt off.

Tuning — the weekly loop

CONTROL	WHEN TO REACH FOR IT
Segment toggles	A whole beat is noise right now — residential storage, EVs, consumer gear. Turn the segment off and the gate zeroes those stories even when they name a company you cover.
Topic weights	You are working a specific problem — interconnection queues, domestic content, fire codes. Weight it up for the duration.
Companies	Synced from Profiler automatically. Adding a dossier there adds it to the scorer here; the two never drift.
Calibrate	A guided rating pass. It deliberately samples mostly mid-scoring articles with some high and some low mixed in, because your judgment on borderline cases carries far more information than your agreement on obvious ones.
Score it	Paste any article and see exactly what the rubric would give it, band by band. The fastest way to understand a surprising ranking.
Source stats	Which outlets are actually earning their slot. Evidence for adding or retiring one.
Presets	Apply a whole tuning profile to an edition at once, rather than flipping 29 toggles by hand.

THE ONE TRAP WORTH KNOWING

Segment vocabularies live in the running spreadsheet, not in the source file — so editing the words in code changes nothing until the seed's version marker is bumped *and* you press **Sync now**. Rebuilding alone does not sync. If you have customised a segment's terms yourself, it is deliberately never overwritten, which also means it will not receive future vocabulary updates.

Getting the most out of it

- **Click through, don't just skim.** The engagement boost is built from real click-throughs, so a week of reading headlines in the email teaches the scorer nothing. Opening the three stories you actually care about is the cheapest tuning available.
- **Thumbs-down is worth more than thumbs-up.** Agreement confirms what the rubric already believed. A rejection is new information, and it is what moves the ranking.
- **Toggle segments before you touch sources.** Almost every "this edition is noisy" problem is a segment problem, not a source problem — and a segment toggle is reversible in one click, where retiring an outlet costs you a beat.
- **Trust the unmarked-geography default.** If you ever find yourself wanting to penalise articles that do not mention the US, remember that most of the coverage you want does not mention it either.
- **Use "Score it" when a ranking surprises you.** Guessing at the rubric wastes more time than checking it. The band-by-band breakdown usually shows a segment you forgot you turned off.
- **Read the footer count on thin days.** "104 scanned, 14 relevant" is a quiet news day. "104 scanned, 3 relevant" with a held-back count of zero is usually a feed that has gone dark — check the background-errors panel.
- **Let editions specialise.** Rather than one edition tuned to a compromise, run BESS and AIDC separately — and add a third for whatever you are actively pitching. Presets make that a one-click setup.
- **Bump the version when you edit segment terms.** This is the single change most likely to look like it worked and silently do nothing.

Where things go wrong, and what the app tells you

SYMPTOM	WHAT TO LOOK AT FIRST
Edition is unusually short	Why thin? — distinguishes a quiet day from a broken feed from an over-tightened tuning profile.
An obviously relevant story is missing	Score it on the article's URL. Usually an off segment zeroing the company match, or a geographic multiplier you did not expect.
The same story appears twice	Clustering did not recognise them as one event. Thumbs-down the weaker one.
Summaries look truncated	The edition hit the storage cap and was trimmed in order. The footer will say so.
A term change did nothing	The seed version was not bumped, or Sync now was not pressed before the rebuild.
A source stopped contributing	Source stats and the background-errors panel. It may have been bot-walled rather than gone quiet.

Reading an edition at a glance

WHAT YOU SEE	WHAT IT MEANS
Amber text	Analysis — the desk's read, not reported fact. The footer key says so on every edition, which is why a figure inside a reported summary is never amber.
Highlighted figures	Concrete numbers pulled out of the prose so you can scan an edition for quantities without reading it.
Section headings	Companies, Market, Incidents — assigned by what the article matched, not by the outlet it came from.
The footer count	Relevant / scanned / held back. The held-back figure survives any trimming, so it is always the true count.

Where the news goes after you read it

An article read once and forgotten is worth little. Scraper's second job is to hold everything it scored in a queryable corpus, and hand the right slice to the other two apps when they need it.

What the corpus holds

Every scored article — including the ones that never made an edition — is stored with the Profiler slugs it matched, its summary and analysis, its event type, and the figures captured from it. A daily pass also mines **edge candidates**: articles naming two covered companies, which may signal a relationship the graph does not yet know about.

TIMELINE	CANDIDATES	READ-ONLY, BOUNDED
One company's news since a given date, newest first, deduplicated — what a dossier refresh reads to find what it missed.	Co-mentions awaiting a judgment about whether a real relationship actually exists between the two.	The route cannot write and returns at most 200 items — a window onto the corpus, not a control surface.

Who reads it, and how it is protected

- **Profiler dossier refreshes.** Every refresh pulls the company's timeline since the dossier's last update plus its pending candidates, verifying each promising item against the article itself before promoting it.
- **Classroom's weekly pipeline.** Corpus items are one source the curriculum run reads when deciding what a seller now needs to know.
- **Access is token-gated.** The route requires a shared secret held server-side in both scripts, never committed and never quoted back. A session without it proceeds without news context and says so — rather than failing, or inventing coverage.
- **The gate follows the content.** Corpus-derived material is contributor-tier wherever it lands — a lesson built from it inherits that gate automatically.

ONE OUTLET IS A LEAD, NEVER A FACT

Nothing crosses the bridge automatically. A single-outlet claim is a lead requiring corroboration; corroborated and company-confirmed items rank higher; every item is verified against the article itself. And there is no write-back — unqualified items are simply left for the daily reconcile pass. Nothing has to stay in sync, so nothing can fall out of sync.

Where Scraper sits

PROFILER — THE SKELETON	SCRAPER — THE BLOOD	CLASSROOM — THE INSTITUTION
Who everyone is. Supplies the covered-company list driving the 40-point company band.	What just happened. Feeds refreshes, relationship candidates and curriculum material back.	What it means and what to retain. Turns both into taught, gated knowledge.

IN ONE LINE

Scraper answers *"what changed today, does it matter to me, and why did you think so?"* — and keeps everything it read, so the answer is still there months later when a dossier needs it.