

Megmeet — Interview Brief (Sales Role)

What they make, where they stand, the objections you carry, and what to ask

Prepared 14 August 2026 · **Source** megmeet.profile.json, profileVersion 2 (updated 2026-08-09) · **Companion** megmeet-lesson-plan.md, which teaches the physics

Sourcing Every claim traces to the dossier; analytical reads are labelled as reads · **Classification** Internal — interview preparation

PRESENTATION STYLE — BloombergNEF Research Report, the canonical Profiler export skin. Facts lead, labelled analysis follows, and the two never blur.

Tuned for: a sales role, interviewing with **Yuan Meng, Head of NA Sales & Marketing**.

Built from `live-site-pages/profiler-data/megmeet.profile.json` (profileVersion 2, updated 2026-08-09).

Companion to `megmeet-lesson-plan.md`, which teaches the underlying physics — this file is the interview-facing layer: products, competitive position, recent activity, hard questions, and questions to ask.

Sourcing discipline. Everything below traces to the dossier. Where the dossier carries a labeled analytical read rather than a sourced fact, this file says so — because asserting an interpretation as fact to someone who works there is the fastest way to lose credibility.

CONTENTS

- | | |
|--|---|
| 01 The one sentence | 09 Where they stand |
| 02 Who you're meeting, and what that changes | 10 The July 2026 FCC inverter rule — what it does and does not touch |
| 03 The five facts that carry the conversation | 11 Recent activity (newest first) |
| 04 The 60-second answer to "what do you know about us?" | 12 The hard questions |
| 05 Company shape | 13 Questions worth asking them |
| 06 The North America picture — the actual job | 14 Vocabulary |
| 07 The five objections you would carry into every meeting | 15 Twelve-question self-test |
| 08 What they make — the product chain is the catalogue | |

The one sentence

Megmeet is a Shenzhen power-electronics manufacturer with six business lines and ¥9.4B of revenue, deliberately burning its profit margin to get designed into the next generation of AI data-center power — the 800 VDC architecture

NVIDIA is driving toward its 2027 megawatt racks. **The old appliance business is paying for the AI bet, and the bet has not paid off yet.**

Who you're meeting, and what that changes

Yuan Meng does not appear anywhere in the dossier — not in `decisionMakers`, not in the source set. Nothing below is a claim about them personally. Do not walk in with a manufactured sense of their background; if it comes up naturally, ask.

What the dossier does tell you is about the **function**, and that is enough:

- **The NA commercial organisation is young.** The only US-facing leader on record is **Roya Movahedi, CMO for US/international** — quoted spokesperson in every 2025–26 English release (NVIDIA HVDC sidecar, OCP, Dallas lab). She is not listed in the China filings, which says the international commercial layer was built recently and separately from the China corporate structure. A dedicated head of NA sales and marketing is consistent with that build-out continuing. (*Reasonable inference from the record, not a stated fact.*)
- **You will be assessed on selling, not physics.** Nobody needs you to design a power shelf. They need to know you can run a **design-in** sale: long qualification cycles, technical buyers, a spec won upstream, and a brand carrying a real objection into every first meeting.
- **Your value is what they cannot hire in Shenzhen.** A mainland manufacturer building a North American business is buying relationships, market access, and objection-handling in a market where its brand is unknown and its country of origin is a live procurement question.

If you have AIDC power/storage pipeline experience, lead with the map, not the résumé. The asset is that you already carry the competitive landscape — who the hyperscalers and neoclouds are, which developers and contractors hold the decision, where power constraints bite, how a component vendor gets specified into a campus. A manager building an NA desk wants someone who arrives with the territory mapped, not someone who spends two quarters learning it.

The five facts that carry the conversation

FACT	NUMBER	WHY IT MATTERS
They are on NVIDIA's list	Only mainland-China power vendor named for GB200 NVL72	In a server-power market ~74% held by Taiwanese suppliers, being the one mainland name NVIDIA called out is the door they walked through
The pivot is expensive	FY2025 net profit -66.6% to ¥146M	Revenue grew 15%; Q4 was a loss; operating cash flow went negative for the first time
AI revenue is not here yet	"Immaterial" through 2025 — <i>their own characterization</i>	Company told 21jingji its AI-server PSU orders were still small-batch, small-value, uncertain timing
But the ramp has started	Power products +67% YoY in Q1 2026 to ¥817M	The bull datapoint most candidates miss; Q1 total revenue +20.4%, all six segments growing
The verdict lands soon	H1 report ~27 Aug 2026	Consensus embeds a ~6x FY2026 profit rebound premised on AI power; H1 is the first hard test

Do not assert: trade press reported Megmeet displacing Lite-On as the #2 NVIDIA power-shelf source. The company has never confirmed it and its own statements point the other way. Say "I've seen it reported, though I don't think you've confirmed it."

The 60-second answer to "what do you know about us?"

SAY IT LIKE THIS

"You're a power-electronics company, not really an appliance company or an AI company — that's the thing people get wrong. The founders came out of Huawei Electric and Emerson Network Power, so conversion is the core competence, and you've pointed it at six industries: appliance controls, EV and rail, industrial automation, welding, precision components, and power products. What's interesting right now is that the power-products group is being turned into an AI data-center business, funded out of the mature lines — R&D has been 11–12% of revenue for a decade, and FY2025 profit took a two-thirds hit to pay for it. The bet is that the shift from AC distribution to 800 VDC reshuffles who supplies rack power before the incumbents lock it down, and being the one mainland vendor NVIDIA named for GB200 is the beachhead. And I know the H1 report at the end of this month is really the first proof point on whether that revenue is showing up."

Why it works: it names the real identity (a conversion-technology platform, not a product company), credits the founding DNA, frames the profit collapse as a choice rather than a failure, and lands on a catalyst date almost nobody else will know.

Company shape

BUSINESS GROUP	FY2025 SHARE	DIRECTION
Smart home appliance controls	37.9% · ¥3.56B	Largest, and the only segment that shrank (–4.8%) on domestic price wars and soft India AC demand. The wallet funding the pivot
Power products (incl. AI data center)	28.5% · ¥2.68B	+14% FY2025, then +67% in Q1 2026 . The strategic centre of gravity
NEV & rail transit	12.2% · ¥1.14B	Fastest growth (+109%) but lowest margin of the majors (15.3%) — volume land-grab, dilutive to mix
Industrial automation	9.3% · ¥877M	+40%. Drives, servos, PLCs. ~28% gross margin
Intelligent equipment (welding, microwave, oil)	6.4% · ¥605M	+31% at 38.5% gross margin — the highest in the company . The quiet profit engine
Precision connection	4.8% · ¥449M	+18% at 5.1% gross margin; company calls it an investment-and-integration period

~40% of revenue earned overseas. 8,894 employees at end-2025 (~3,100 in R&D). Ten R&D centres, eight manufacturing bases across China, Thailand, India, the US, Germany. Founded 2003; SZSE-listed 2017; HK A+H dual listing filed June 2026.

The North America picture — the actual job

The single most important nuance

Megmeet's expected **first hyperscaler volume delivery is domestic — a Chinese hyperscaler, in 2026**. That means **North America is genuinely greenfield**: the hardest market, the strongest incumbent, the most sceptical procurement, and no reference win yet to point at. Whoever takes this desk builds the first NA proof case rather than managing an existing book.

Most candidates read "first hyperscaler volume delivery expected 2026" and assume a US buyer. Catching that it says *domestic* shows you read the source rather than the summary — and it reframes the role honestly. Frame it as what it is: a build, and the reason the job exists.

What you have to sell with

ASSET	WHY IT MATTERS COMMERCIALY
The NVIDIA ecosystem seat	Only mainland-China power vendor NVIDIA named for GB200 NVL72, plus a co-developed 800 VDC sidecar shown publicly at OCP. The spec is won inside reference designs before an RFP exists — that seat is the door, and it is the one thing a competitor cannot simply price against
The Dallas lab	June 2026, 360 kW active, 1.5 MW roadmap. A sales instrument, not a facility announcement — validation on US soil shortens qualification cycles, lets customer engineers touch hardware without a Shenzhen trip, and answers "can you support us here?" with a place rather than a promise
Existing Western OEM relationships	Named network-power customers include Ericsson, Cisco, Juniper, Arista, Accton — Western tier-ones that have already run the vendor-qualification gauntlet. The strongest available counter to "we've never heard of you"
Full-chain capability	Grid-to-GPU rather than one box: SST, sidecar, shelves, BBU/supercap, PSUs, point-of-load. Lets you sell an architecture conversation instead of a line item
Spec leadership	Ruby-grade 5.5 kW M-CRPS at >100 W/in ³ ; the jump from 33 kW air to 110 kW liquid. In a standardised form-factor market, spec is the entry ticket — and standardisation is how a challenger takes a socket
Financed capacity	The ¥2.66B placement (Feb 2026, chairman subscribing personally) and the pending HK listing fund Thailand phase II and the rest of the build-out. If a customer asks whether you can scale with them, the capital is identifiable

The sales motion you would run

Per the dossier, power products go to market through **direct and ODM design-ins with server OEMs, cloud providers, and the NVIDIA reference ecosystem**; network power sells direct to telecom/networking OEMs. In practice:

- **Design-in selling, not transactional.** Long cycles, engineering buyers, qualification gates, a win that locks for a product generation. Revenue arrives quarters after the work
- **The spec is won upstream.** Reference designs and the partner slate decide who gets to bid; by the time procurement opens the outcome is often set
- **The buying centre is split.** Rack architecture, thermal, supply-chain risk, and cost are different people. The China-origin objection lives with supply-chain risk, not the architect — a different conversation with different evidence
- **ODMs are both channel and gatekeeper.** A hyperscaler may set the spec while an ODM holds the PO. Knowing which door an account opens is the qualifying question

Country of origin — handle it precisely

The dossier's read: the geographic build-out (Thailand phase II, India localisation, the Dallas lab, the HK listing) is **supply-chain de-risking aimed at US hyperscaler procurement** — a mainland company packaging itself to be buyable by American AI-infrastructure buyers despite tariff and geopolitical friction. (*Labeled analytical read, moderate confidence — not company guidance.*)

Do not overclaim the regulatory picture. The restricted-entity and tariff machinery from the battery world — 1260H listings, stepped tariffs on Chinese cells, the draft inverter rule — is documented for **battery and inverter suppliers, not for Megmeet**, which is not recorded as being on those lists. Importing that framing wholesale describes someone else's problem and an insider will know. The safe, honest framing: country of origin is a live procurement question for any mainland supplier into US AI infrastructure, and the company's answer is manufacturing and validation footprint outside China.

The five objections you would carry into every meeting

A sales interview is really an audition for objection handling. The strongest move is to raise these yourself, unprompted, and show you already have answers.

WHAT THEY SAY	WHAT THEY'RE REALLY ASKING	HOW YOU ANSWER
"You're a Chinese supplier."	Will procurement or legal block this later and waste my qualification cycle?	Answer with footprint, not reassurance: manufacturing across Thailand, India, the US, Germany; Thailand phase II funded; a Dallas validation lab on US soil; an HK listing in progress. Then the proof that lands hardest — Ericsson, Cisco, Juniper, Arista and Accton already buy Megmeet power and have already run their supplier audits. Never promise regulatory outcomes
"We've never deployed you at this scale."	Am I the guinea pig?	Concede the specific, redirect to the general: new to hyperscale AI racks, not new to high-volume power manufacturing — global top-three in OEM/embedded power supply sales for 2024, ~8,900 employees, ~3,100 in R&D, ¥9.4B revenue. Then the structural answer: standardised form factors let a buyer take a limited, low-risk slot first
"Your financials look shaky."	Will you still be here in three years, mid-programme?	Don't flinch — it's public and they've read it. Profit fell two-thirds because the company chose to spend ahead of revenue; capex is financed by a ¥2.66B placement the chairman subscribed to personally plus a pending HK listing; Q1 2026 had all six segments growing with power products +67%. The diversified base is why the AI bet doesn't have to work on the first try
"Delta already has this socket."	Why re-open a qualification that works?	Never attack Delta — they're good and the buyer knows it. Sell the second source: ~74% geographic concentration is a risk they already manage elsewhere in the BOM. Then sell timing — the architecture change forces re-qualification anyway, so the switching cost you're asking them to absorb is one they're already absorbing
"What if 800 VDC slips?"	Am I buying a roadmap that may not arrive?	Point at the bridge. The sidecar exists because buildings on 480/380 VAC need megawatt racks <i>now</i> ; the near-term catalogue (33 kW and 110 kW shelves, Ruby M-CRPS) sells into today's architecture regardless. The roadmap is upside, not the thesis

The line that ties it together: "The way I'd think about this territory is that we're not asking anyone to replace Delta. We're asking to be the qualified second source during a re-qualification they have to do anyway — and the Dallas lab exists so that costs them a lot less time than it sounds like."

What they make — the product chain is the catalogue

You are not being hired to design this — you are being hired to sell it to people who do. The bar is being credible with an engineer for twenty minutes. The trick: Megmeet's AI catalogue *is* the power chain, box by box.

STAGE	PRODUCT	WHAT TO KNOW
10–35 kV AC	The grid arrives (medium voltage)	High voltage because low voltage means enormous current, and losses rise with the <i>square</i> of current
→ 800 V DC	Solid-state transformer (SST)	A transformer built from switching electronics instead of just an iron core. MV grid AC → 800 VDC in one system at >98.5% , megawatt scale. Replaces the step-down transformer, the rectifier stage, part of the central UPS. The product that would make them an infrastructure vendor rather than a component vendor
480/380 V AC → 800 V	800 V HVDC sidecar	The bridge shipping today. A cabinet standing <i>beside</i> the rack converting building AC to 800 VDC, feeding >1 MW into a single rack without rebuilding the facility. Claims vs legacy in-rack AC: +5% efficiency, –45% copper, up to –30% TCO . Announced for NVIDIA Kyber May 2025 — first mover among mainland vendors
800 V DC bus	BBU + supercapacitor shelves	The UPS dissolved into the rack. BBU (16.5 kW) rides through seconds-to-minutes; supercap shelf (15 kW) absorbs <i>millisecond</i> swings. Both are needed because GPU training loads slam between compute and communication phases
33 kW / 110 kW	Power shelves	Slide-in 19-inch trays feeding the rack DC busbar. 33 kW air-cooled (1U, hot-swap, N+1/N+N) and 110 kW liquid-cooled for GB300/Vera Rubin racks at 45 °C warm-liquid inlet. The jump to 110 kW liquid (May 2026) is the clearest signal they track NVIDIA's roadmap rather than follow it
800 W – 5.5 kW	CRPS / M-CRPS server PSUs	Standard form factor — any compliant module fits any compliant chassis, which is exactly how a challenger gets in , because sockets are re-competed on spec. Ruby-grade M-CRPS launched Mar 2026 at 5.5 kW, >100 W/in³
→ ~1 V DC	Power brick	Point-of-load converter next to the GPU, dropping the bus to the ~1 V at hundreds of amps the silicon drinks

The sentence that shows you get it: "The reason 800 VDC matters isn't efficiency for its own sake — it's that a classic rack drew 5–15 kW and an NVIDIA-class rack is heading to a megawatt. At 415 VAC a megawatt is ~1,400 A, which means busbar like railroad track. The old chain physically stops scaling, so you convert once, early, at high voltage, and distribute DC instead."

The connective insight for the other segments: every other business line is the *same physics* aimed at a different industry. A variable-frequency drive rectifies grid AC to DC then re-synthesizes AC at a chosen frequency (motor speed follows frequency; motors are ~45% of global electricity). Appliance controls are the same inverter sized for a compressor. EV on-board chargers and DC-DC units are the same stages in a vehicle. Welding is precision current control — and their highest-margin line.

Where they stand

PLAYER	POSITION	WHAT TO KNOW
Delta Electronics	~41% share, Taiwan	The incumbent to beat. Deepest published 800 VDC stack, first-qualified across GB200/GB300 shelf sets, ~70% of Blackwell-platform PSUs

PLAYER	POSITION	WHAT TO KNOW
Lite-On	#2, ~30%, Taiwan	BBU franchise is its structural edge — BBUs went from optional (GB200) to standard (GB300). Slightly behind Delta on the 800 V clock
Taiwanese bloc	~74% of server power	Delta, Lite-On, AcBel, Chicony, Advanced Energy. This concentration is why a mainland challenger interests NVIDIA at all
Megmeet	Only mainland vendor NVIDIA named	Competing on architecture timing, not price . Co-developed 800 VDC sidecar shown at OCP Global Summit Oct 2025

****The strategic argument (labeled analysis — moderate confidence in the dossier, *not* company guidance):**** when a technology generation turns over, the incumbent's qualifications, installed base and tooling partially reset, and every socket is re-qualified against a new spec. Megmeet's wager is that it gets designed in during that reset window, before incumbents re-consolidate around 800 VDC as they did around AC. Present this as "my read is...", never as something Megmeet said.

The second half of the position: Thailand phase II capacity, India localisation, the Dallas test lab, and the Hong Kong listing all point one way — a mainland company packaging itself to be procurable by American AI-infrastructure buyers despite tariff and geopolitical friction. Named network-power customers already include Ericsson, Cisco, Juniper, Arista, Accton, so the Western-OEM muscle exists and is being redirected.

The July 2026 FCC inverter rule — what it does and does not touch

On **28 July 2026** the FCC added **foreign-produced connected power inverters** (and advanced robotic devices) to its Covered List, effective immediately. Covered-List products cannot receive FCC equipment authorization, which bars them from being imported, marketed or sold in the US.

Note on sourcing: this section was verified against current sources on 2026-08-14, *not* from the dossier — the Huawei dossier still records this as a **draft, China-specific rule** from the 30 June Reuters report and is stale. See the note at the end of this section.

The three things almost everyone gets wrong

- 1. It is not a China rule.** A device is "foreign-produced" if it fails the Buy American *domestic end product* test at 48 CFR § 25.101(a) — manufactured in the US with domestic component cost above 65% through 2028, rising to 75% in 2029. It turns on **where a product is made, not who makes it**. An American brand manufacturing offshore is caught. **So are Delta and Lite-On.**
- 2. It is prospective only.** Models authorized before 28 July keep their authorization and can still be imported and sold. This is a new-authorization freeze, not a market withdrawal.
- 3. There is a way out.** Conditional Approval from the Department of War or DHS, applications until 1 January 2028, requiring corporate-ownership disclosure, a full bill of materials and supply-chain map, and a **time-bound US onshoring plan** with capital commitments and milestones. No published review timeline, so approval is not guaranteed.

The definition is a two-part test — both prongs must be met

PRONG	WHAT IT SAYS
Conversion	A bi-directional power device or system converting DC → AC or AC → DC — enumerated as microinverters, string inverters, central inverters, and hybrid/battery-based inverters
Connectivity	Contains components enabling remote communication, control, sensing, data collection or monitoring through Wi-Fi, cellular, Bluetooth or another similar connection . Stated rationale: remote firmware push by a foreign adversary

Does it hit Megmeet? Three different answers

PRODUCT LINE	EXPOSURE	REASONING
AI data-center power (PSUs, shelves, sidcar, SST)	Probably out — but unsettled	These convert AC → DC, which is inside the conversion prong. But they are unidirectional , not bi-directional; none is a micro/string/central/hybrid inverter; and management is wired (PMBus on the PSU, IPMI/Redfish/SNMP over Ethernet on the shelf), not Wi-Fi/cellular/Bluetooth. On a plain reading of both prongs they fall outside. The live ambiguity is " another similar connection " — the industry is already flagging that wired-connected <i>central</i> inverters sit in a grey zone, and the FCC has told manufacturers not to assume devices are out of scope
PV / storage / EV-charging components	Genuinely exposed	Where Megmeet actually intersects the rule. Hybrid and battery-based inverters are named categories, storage PCS is bi-directional by design, and EV charging equipment is within the reported scope. Note the timing: Megmeet launched a 40 kW AC-DC and DC-DC EV-charging module platform on 31 July 2026 — three days after the rule took effect
Everything else (appliance, industrial, welding, on-vehicle)	Out	Not power inverters within the definition; vehicles are among the excluded systems

The answer that will land

SAY IT LIKE THIS

"My read is it mostly misses the rack-power business. The definition is a two-part test — a bi-directional inverter and wireless connectivity — and a server PSU or power shelf is unidirectional AC-to-DC with wired management, so on a plain reading it's outside. Where it does bite is the PV, storage and EV-charging side, which is a much smaller slice. The part I find genuinely interesting is that it's an origin test, not a China test — Buy American domestic-end-product — which means Delta and Lite-On are foreign-produced too. If the interpretation ever stretches toward IT power, it doesn't hit us asymmetrically; it hits the incumbents who have far more US SKUs to re-authorize. And the conditional-approval route runs through a US onshoring commitment, which is a strategic question rather than a compliance one."

Where to hedge

No source found addresses data-center or IT power supplies directly — **the read above is inference from the definition, not a sourced ruling**, and the FCC's own FAQ and determination PDF returned 403 and could not be read. Two further variables turn on facts not public: whether any specific Megmeet SKU carries a wireless module, and where each is manufactured. Say "my read", not "the rule says". If pushed hard, the correct answer is that this is a question for counsel and the certification body, not a salesperson — that answer is a strength, not a dodge.

Dossier staleness flagged: `huawei-digital-power.profile.json` records the 2026-06-30 Reuters draft-rule report and a strategyRead judging the Megmeet spillover "unconfirmed and scope-dependent". Both predate the 28 July action

and the country-agnostic scope. The Huawei dossier and `AIDC-MARKET-REPORT.md` §6.4/§8.4 should be refreshed.

Recent activity (newest first)

The pattern to notice: **the product and capacity news is relentlessly positive while the financial news is relentlessly hard** — and both are true at once. That tension *is* the company right now.

DATE	EVENT	READ
2026-07-20	Three 10% limit-downs in four sessions; ~-31% in 20 days off the June peak	Sell-side cited profit collapse, negative operating cash flow, >¥6B guarantee quota, HK dilution risk. The AI premium is real but fragile
2026-06-29	HK A+H dual listing filed (Huatai International, Citigroup)	Offshore raise to fund AI-power capex; arguably necessary given negative operating cash flow, dilution the accepted price
2026-06-26	Dallas, Texas AI-DC power test lab opened — 360 kW active, 1.5 MW roadmap	Validation capacity next to US hyperscaler customers — localisation beyond "made in Thailand"
2026-05-29	NVIDIA MGX: 110 kW liquid-cooled power shelf , 800 VDC rack conversion, BBU/supercap shelves	Moving up from 33 kW air to the liquid-cooled class for GB300 / Vera Rubin racks
2026-04-28	FY2025 results + Q1 2026	FY2025: revenue ¥9.40B (+15.1%), net profit ¥146M (-66.6%), ex-one-off -93%, Q4 a loss, OCF -¥139M. Q1 2026: revenue ¥2.788B (+20.4%), all six segments growing, power products +67% to ¥817M
2026-04-13	Full 800 VDC grid-to-GPU chain shown at Open AI Infra Summit, Beijing	SiC/GaN front end, sidecar, power/BBU shelves, digital-twin control — full-chain vendor, not a single-shelf supplier
2026-03-03	Ruby-efficiency M-CRPS 5.5 kW server PSU	Past Titanium to Ruby at >100 W/in ³ — spec leadership is the entry ticket to AI-server sockets
2026-02-06	¥2.66B private placement completed at ¥85.01/share	Subscribers included E Fund and UBS; the chairman participated personally . Funds Changsha R&D, Thailand phase II, Zhuzhou phase III — the capex is financed
2026-01-28	FY2025 profit warning: net profit -66% to -72%	Blamed on AI-power R&D, overseas build-outs, price wars, FX losses — the pivot's cost landing before its revenue
2025-11-04	Company said AI-server PSU orders were still small-batch , small-value, uncertain timing	Their own words to 21jingji, and they contradict market share-gain reporting. The single most useful thing to know
2025-10-13	Co-developed 800 VDC sidecar shown at OCP Global Summit	Public NVIDIA co-development visibility — the ecosystem-credibility milestone
2025-05-21	800 V HVDC sidecar for NVIDIA Kyber announced (>1 MW/rack)	First mover among mainland vendors. Where the AI story starts

Market context, if raised: ~¥24 (early 2024) → ¥190 intraday peak / ~¥94.7B market cap (3 Jun 2026) → ¥108.71 / ~¥63.6B (29 Jul 2026). Handle lightly — know the arc, don't volunteer an opinion on your prospective employer's share price.

The hard questions

In each case the winning move is the same: **acknowledge the real difficulty, then show you understand why the company chose it.** Do not be a cheerleader — people who work there know the numbers better than you do.

"What do you make of our FY2025 results?" — Name the trade honestly: revenue +15%, profit –two-thirds, because R&D and overseas capacity were pulled forward ahead of the revenue they serve, with FX losses and price competition on top. Then make the analytical point: the useful question isn't whether profit fell, it's *whether the spending bought position*. Evidence it did — NVIDIA ecosystem membership, the co-developed sidecar, the jump to 110 kW liquid, power products +67% in Q1. Evidence still outstanding — revenue at scale, which H1 will show.

"Is the AI data-center business real, or is it a story?" — The honest answer earns more than the enthusiastic one. Per the company's own late-2025 statements, AI orders were small-batch and immaterial, most projects pre-mass-production. What has changed is directional: batch deliveries of custom server power began end-2025, first domestic hyperscaler volume delivery expected 2026, power products +67% in Q1 2026. So: real and starting, not yet material. Then the sharp bit — consensus prices a ~6x FY2026 profit rebound, so the market has already decided; H1 shows whether it was right.

"Why would a hyperscaler buy from us instead of Delta?" — Three arguments together. *Architecture timing*: during a reset, sockets are re-qualified and the incumbent's lead partially resets. *Supply-base diversity*: a ~74% Taiwanese market is a concentration risk large buyers actively manage. *Full-chain capability*: Megmeet can quote grid-to-GPU rather than one box, which is rare. Then be candid — Delta is first-qualified across the shelf sets and holds most Blackwell PSU volume, so this is a wedge, not a coronation.

"What worries you about us?" — A real answer, calmly delivered, is a strength. The expectations-vs-disclosure gap is the honest one: the company said AI revenue was immaterial and consensus is priced for a 6x rebound — a lot of load on one half-year print. Second, the fastest-growing segment (EV/rail, +109%) runs at 15.3% gross margin, so growth is dilutive to mix. Third, negative operating cash flow means the build-out depends on capital markets, which is presumably part of why the HK listing exists. Frame these as things you'd want to understand better, not accusations.

"How would you open a hyperscaler or ODM account for us, cold?" — The wrong answer is a pitch; the right answer is a sequence. Start where the spec is set, not where the PO is cut: the NVIDIA ecosystem seat and the reference-design layer are why a first meeting happens at all. Qualify which door the account opens — does this buyer set the rack architecture while an ODM holds the PO, or do they buy direct? Find the person who owns supply-chain risk, because that is where the real objection lives. Make the ask small: a standardised form factor lets you compete for one limited slot rather than the whole socket. And use the Dallas lab early — "come touch it in Texas" converts scepticism faster than any deck.

"What would your first ninety days look like?" — Concrete and modest. Map which accounts are reachable through the existing Western OEM relationships versus genuinely cold. Get fluent enough on the chain to hold a technical call without pulling in Shenzhen. Understand what the Dallas lab can and cannot validate at 360 kW today, since that bounds what you can credibly promise. Find where qualification cycles have stalled before — in design-in selling the losses teach more than the wins. Then say the honest thing: with cycles this long the first ninety days builds qualified pipeline, not revenue, and you would want to agree upfront which leading indicators you will both judge it by.

Say / don't say

SAY

"My read is the profit collapse was a deliberate purchase of position — and H1 is the first real scoreboard."

DON'T SAY

"The profit drop isn't a concern." It plainly is; pretending otherwise signals you didn't read the filings

SAY	DON'T SAY
"You're the only mainland vendor NVIDIA named for GB200 NVL72 power."	"You've taken the #2 shelf slot from Lite-On." Reported, never confirmed, contradicted by the company
"The 800 V transition is a vendor-list reset on a published clock — first shipments this year, Kyber racks 2027."	"800 volts is just more efficient." True but shallow — the driver is that megawatt racks make AC distribution unworkable

Questions worth asking them

With a sales leader the bar is specific: ask about the **go-to-market**, not the company. Pick three or four; don't run the list.

About the territory

1. **"Am I building the first North American reference win, or is there one already that I'd be scaling?"** — The most important question, because the answer defines the job. You already suspect the former (first hyperscaler volume delivery is expected to be domestic), so asking shows you read the record and are walking in clear-eyed.
2. **"Where does the North American number actually come from first — hyperscalers directly, the ODMs, or growing the network-power base at accounts like Cisco and Arista?"** — Names the three plausible paths and makes them tell you which is real.
3. **"Does the ecosystem seat translate into ODM design-ins, or do you still win each one separately?"** — The real commercial question underneath the thesis.

About the sale itself

1. **"When a US customer raises country of origin, what's the answer that's actually been working?"** — The best question available. Respectful, assumes competence, gets them talking about what they think about most, and whatever they say is the most useful thing you could learn before day one.
2. **"Is the Dallas lab a sales instrument or an engineering one right now, and who walks customers through it?"** — Treats their newest investment as a commercial asset, which is how a sales leader thinks about it.
3. **"How is the desk structured — do I own accounts end to end, or is it a spec-in role that hands off to whoever holds the ODM relationship?"** — Practical, and shows you know design-in selling has that seam.
4. **"What does pre-sales engineering support look like in-region?"** — In a technical design-in sale this is the difference between a workable territory and an impossible one.
5. **"Has the July FCC covered-list action changed anything about how you're planning US authorizations for the power-products line — and does it move the calculus on US manufacturing versus Thailand?"** — Current, specific, and it shows you understand the rule is an origin test with a US-onshoring escape hatch rather than a China ban.

The one that will be remembered: *"When the half-year numbers land at the end of the month — does that change anything about how you sell here, or is it mostly a capital-markets event from where you sit?"* It proves you know the calendar, it is warm rather than adversarial, and it draws a line between the equity story and the sales motion that most candidates never see.

Vocabulary

TERM	MEANING
800 VDC / HVDC	High-voltage DC distribution inside the data hall — the architecture transition everything revolves around
SST	Solid-state transformer — MV grid AC → 800 VDC in one electronic system, >98.5%
Sidcar	Cabinet beside the rack converting building AC to 800 VDC — the retrofit path to megawatt racks
CRPS / M-CRPS	Common Redundant Power Supply — the standard server-PSU form factor; "M" is the higher-power variant
BBU	Battery backup unit — seconds-to-minutes ride-through on the rack DC bus
Supercapacitor	Far less energy than a battery, delivered far faster — catches millisecond load swings
N+1 / N+N	Redundancy: one spare module beyond the load, or full duplication
Hot-swap	Replacing a failed module without powering anything down
PUE	Power usage effectiveness — facility power ÷ IT load. 1.5 ordinary, <1.1 the ambition. The industry scoreboard
Titanium / Ruby	Server PSU efficiency grades; Titanium ≈96%, Ruby the newer top tier
GB200 / GB300	NVIDIA rack-scale GPU systems; NVL72 = 72 GPUs joined as one unit
Kyber / Vera Rubin	NVIDIA's next rack generation and platform — the racks that need 800 VDC, arriving 2027
MGX	NVIDIA's modular server reference architecture partners build to
OCP	Open Compute Project — the open hardware standards body whose summit is where this ecosystem shows its work
SiC / GaN	Silicon carbide, gallium nitride — wide-bandgap semiconductors that switch faster and hotter than silicon, enabling compact high-efficiency conversion
VFD	Variable-frequency drive — the industrial-automation product; controls motor speed by synthesizing AC at a chosen frequency
PCS	Power conversion system — the inverter block in battery storage; the bridge to the storage world

Twelve-question self-test

Answer out loud before checking. Recall is what makes it stick.

1. Why does an AI rack need 800 VDC instead of 415 VAC?
2. What does a solid-state transformer replace?
3. Why does the rack need both a BBU and a supercapacitor shelf?
4. What is the commercial difference between the sidcar and the SST?
5. Why is a standardized PSU form factor good news for a challenger?
6. What share of server power do Taiwanese suppliers hold, and who leads?
7. What did Megmeet itself say about its AI order book as of late 2025?
8. What happened to FY2025 profit, and why?

9. What is the single most encouraging recent number?
10. What happens at the end of August, and why does it matter?
11. Which market is the first hyperscaler volume delivery expected in, and why does that matter to you?
12. Name the strongest single answer to "you're a Chinese supplier".

ANSWERS

1. A megawatt at 415 VAC is ~1,400 A — impractical busbar and severe resistive losses; losses scale with the square of current. Converting once at high voltage and distributing DC also removes stages and lets batteries connect natively.
2. The conventional step-down transformer, the rectification stage, and part of the central UPS — collapsed into one system at >98.5%.
3. Different timescales: supercaps absorb millisecond transients from GPU compute/communication phase swings; batteries carry seconds-to-minutes ride-through.
4. Sidecar is the retrofit product for buildings distributing 480/380 VAC today; the SST is for purpose-built DC campuses and arrives later. Sidecars dominate through ~2027.
5. Any compliant module fits any compliant chassis, so sockets are re-competed on spec rather than locked to an incumbent relationship.
6. ~74%, led by Delta at ~41%, with Lite-On, AcBel, Chicony, Advanced Energy behind.
7. That orders were still small-batch, small-value, uncertain timing, with most projects pre-mass-production.
8. Net profit -66.6% to ¥146M on revenue +15.1% to ¥9.40B; ex-one-off -93%; OCF negative for the first time. Causes: AI-power R&D, overseas build-outs, price competition, FX losses.
9. Power products +67% YoY in Q1 2026 to ¥817M, inside total revenue +20.4% with all six segments growing.
10. The H1 2026 report (~27 Aug). Consensus embeds a ~6x FY2026 profit rebound premised on AI data-center power — the first hard test of whether it is materializing.

Megmeet — Interview Brief (Sales Role) · Internal, interview preparation · Rendered from megmeet-interview-brief.md, which remains the source of truth · Developed by: LightAISolutions