

Zhonhen Electric — US AIDC Market-Entry Strategy Report

Positioning among SST/MV providers, the ride-through and flicker sale, the two relationships, and the neocloud targets

Prepared 22 August 2026 · **Sources** zhonhen.profile.json (41 public sources) · the company deck · fresh web research, cited inline · **Companions** zhonhen-interview-brief.md · zhonhen-deck-summary.md

Handling Contains confidential relationship intel relayed by the company — do not circulate · **Classification** Internal — strategy preparation

SOURCING DISCIPLINE — crux market facts carry numbered citations; first-hand intel from the company is labelled as such and never blended with sourced fact; strategy synthesis is labelled analysis.

Prepared 2026-08-22 · Written from Zhonhen's perspective for the developer's second-round conversation with Zhu Yikun ("Jacky Zhu") · Companions: [zhonhen-interview-brief.md](#) (the reference), [zhonhen-deck-summary.md](#) (their pitch), [zhonhen-lesson-plan.md](#) (the physics)

Handling note. This report contains first-hand relationship intel relayed from Zhu Yikun — the Schneider Electric relationship is **confidential** and must never be repeated outside conversations with Zhonhen itself; the CATL transaction and the neocloud targeting intent are public or verifiable. This document stays in non-deployed study-prep/ and is never cited into the public dossier.

Update (2026-08-22) — the NVIDIA August 2026 white paper is now in hand. Committed at [repository-information/industry-guidance/sources/nvidia-800vdc-white-paper-2026-08.pdf](#) with a full analysis ([nvidia-800vdc-analysis.md](#)) and an in-app study module (Profiler → Industry Guidance). Two findings amend this report: **(a) the validation quote is verified** — p22 names the "Panama Architecture" as one of three canonical TRU implementations, "offering a simpler and highly familiar design philosophy"; **(b) the window tightens** — the 2029 date attaches to *next-gen* SST specifically, while TRU-based ~4.8 MW power blocks (Option C initial implementation, containerized) are being specified now and Option B Power Centers deploy as soon as Q3 2027. The defensible claim is no longer "nothing Western ships until 2029" but "the reference document's near-term block is a TRU — Zhonhen's device class since 2019 — with OEM competition converging sooner." Individual amendments are marked [NV-0826v] below.

CONTENTS

- 01** 1. Key judgments
- 02** 2. The two specs that now sell MV solutions — ride-through and flicker
- 03** 3. The field — where Zhonhen sits among SST and MV-solution providers
- 04** 4. The two relationships — how to spend them
- 05** 5. The neocloud thesis — verified, refined, and named

1. Key judgments

1. **(High confidence) The sellable window is 2026–2029, and it is real.** NVIDIA's own productization ladder — power racks 2H 2026, row power centers 2027, MV-to-DC data-hall power blocks (including SST) as a 2029 target [NV-0826] — reaches Zhonhen's core capability *last*. One-stage MV-to-DC conversion with a delivered fleet is something no Western vendor ships today; the strategy is to monetize that gap, not to out-roadmap the SST future.
2. **(High confidence) Regulation just wrote Zhonhen's US sales narrative.** ERCOT's LCL ride-through rule became binding on 1 August 2026, FERC ordered national computational-load standards on 16 July 2026, and ERCOT has a numerical flicker limit in the pipeline [ERCOT-1308, FERC-RD26, ERCOT-LLWG]. An architecture with batteries on the DC busbar and storage at four depths is *compliance by construction* — the pitch is a compliance artifact, not a brochure (Section 2).
3. **(High confidence) The neocloud thesis is correct, with a refinement that changes the call list.** Owner-operator self-builders — IREN and Crusoe above all — buy MV equipment directly, at speed, with open vendor slots; tenant neoclouds (CoreWeave, Lambda, Nebius) buy only the rack-side chain; xAI fits technically and is closed politically (Section 5).
4. **(High confidence) The Schneider relationship is the highest-value asset Zhonhen holds for the US — worth more than any direct sale.** Schneider's published 800VDC position needs a TRU it does not make, its prefab plants and \$700M US expansion can wrap Zhonhen's blocks in an American brand, UL pathway and service network, and the Vertiv/Hopewind precedent proves the model [SE-800V, CNESA] (Section 4.2).
5. **(Moderate confidence) CATL is a US asset only when kept behind the architecture.** The 49% holdco investment funds the balance-sheet credibility a pilot program needs and completes the source-to-rack power-quality machine — but CATL content triggers the FEOC/1260H stack, so every US product must stay battery-agnostic and CATL stays out of the lead story [CATL-DEF, FEOC] (Section 4.1).
6. **(Moderate confidence) Honesty about the gaps is itself the credibility play.** Zhonhen is on neither NVIDIA 800VDC roster, has no US entity, UL listing or reference site, and its best validation quote — once gated — is now verified: the August 2026 white paper names the Panama Architecture at p22 [NV-0826v]. The wedge is regulatory physics plus schedule — not ecosystem membership — and saying so precisely is what a serious counterparty (and Jacky) will respect.

2. The two specs that now sell MV solutions — ride-through and flicker

Jacky's framing is not just right — it understates how fast the ground moved. In the three weeks before this report, voltage ride-through went from a preference to **binding ERCOT law**, and the flicker problem acquired its first proposed numerical limit. This section is the technical narrative the US sales motion should be built on.

2.1 Ride-through: from incident reports to a rule with teeth

The record that forced the rule. On 10 July 2024, a single failed lightning arrestor near Fairfax, VA produced six successive sub-100ms voltage sags — and ~1,500 MW of load, "exclusively data center-type load," disconnected itself; ~1,260 MW stayed off **for hours**, because UPS disturbance-counting schemes (three sags in a minute → stay on backup until *manual* reconnection) latched the fleet off the grid [NERC-0125]. On 22 July 2026 it happened again, bigger: a transmission fault near Ashburn dropped **over 3 GW** of data-center load off PJM at once [PJM-0726]. ERCOT's

own ledger: **26 events of >100 MW indirect load loss** from computational facilities between January 2023 and September 2025 — facilities shedding 17–95% of consumption on sags as shallow as 0.844 p.u., some taking hours to return [NERC-0126]. ERCOT's VP of System Operations put the systemic risk plainly: sudden frequency changes from mass load trips "can cause other loads and generators to trip offline, potentially resulting in cascading outages" [ERCOT-BOD].

The rule. NOGRR282 + NPRR1308 — approved by the PUCT on 9 July 2026, **effective 1 August 2026** — create the "Large Computational Load" (LCL) class: any site ≥ 75 MW aggregate peak with $\geq 50\%$ computational load. LCLs must ride through defined voltage/frequency envelopes (e.g., sags to 0.50–0.80 p.u. for ≥ 0.5 s; to 0.20–0.50 p.u. for ≥ 0.25 s), **return to $\geq 90\%$ of pre-disturbance consumption within 2 seconds** of voltage recovery, cap disturbance current draw at 125% of normal, use filtered measurements with time delays — and **disturbance-counting trip schemes are explicitly banned**. Failures trigger a root-cause/corrective-plan regime, and ERCOT may order a non-compliant LCL disconnected [ERCOT-1308, ERCOT-SPWG]. Two nuances that matter commercially: the rule is **not retroactive** (loads approved before 14 Nov 2025 are exempt — so the retrofit sale there is voluntary, but every new build is in scope), and facilities energizing **after 1 Jan 2028** face a stricter keep-consuming tier during 0.5–0.8 p.u. sags [ERCOT-SPWG]. Texas SB6 adds the political layer: mandatory remote-disconnection equipment for new ≥ 75 MW loads and ERCOT curtailment authority [SB6].

And it is going national. NERC escalated from a Level 2 alert (Sept 2025) to a rare **Level 3 Essential Action alert (4 May 2026)** with seven required actions on computational loads [NERC-ALERT]; on **16 July 2026 FERC ordered NERC to file mandatory reliability standards for computational loads** — Phase I due 31 December 2026 [FERC-RD26]; MISO's large-load initiative explicitly lists ride-through, ramping and stability criteria [MISO-LL]. A vendor selling ride-through performance in ERCOT today is selling next year's national requirement.

Why conventional chains trip (the physics, all from NERC's incident forensics [NERC-0125, NERC-0126]): static UPS transfer to battery and count disturbances; DRUPS return to grid only manually; constant-power PSUs draw *more* current as voltage falls (0.8 p.u. voltage $\rightarrow$ 1.25 p.u. current; 150–200% spikes observed), tripping overcurrent and neutral-overcurrent protection; cooling VFDs with undervoltage thresholds observed as high as 0.96 p.u. drop out in 50 ms sags and force thermal shutdowns.

Why Zhonhen's architecture rides through by construction (company-stated specs; the mapping is analysis). In the HVDC/Panama chain the battery strings hang **directly on the DC output busbar** — there is no transfer event, no counting logic, and no static switch to latch: a grid sag is buffered at the rectifier and the IT load never sees it [ZH-DOSSIER, DECK]. Centralized MV rectification puts one controlled current path where a conventional facility has hundreds of distributed breakers, letting the 125%-current rule be engineered at a single point; the deck's "no power surge / –60% server-PSU failure" claim is the same property stated as reliability [DECK]. **The commercial translation: make the NOGRR282 envelope the pilot's acceptance test.** Demonstrating a compliant sag ride-through on a container-level test converts a Chinese field record into an ERCOT compliance artifact — the single most persuasive object a US buyer can be handed. (Precision for the call: quote ERCOT's record as "26 events over 100 MW" and attribute the 1.5 GW and 3 GW events to Northern Virginia/PJM — getting the geography right signals command of the primary documents.)

2.2 Flicker: the load's own misbehavior, about to be regulated

The record. Microsoft, OpenAI and NVIDIA jointly published the defining paper (Aug 2025): synchronized training swings aggregate power "by tens or hundreds of megawatts," with the swing spectrum concentrated at **0.2–3 Hz** — **overlapping the torsional and inter-area oscillation bands of grid turbine-generators**, which the authors warn "can cause physical damage to the power grid infrastructure" [ARXIV-PS]. Meta disclosed tens-of-MW instantaneous fluctuations training Llama 3 [LLAMA]. NERC's task force published high-speed recordings of a 50 MW block of an AI

data center ramping at **1.9 p.u. per second** [NERC-LLTF]. ERCOT has briefed since July 2025 on LCL power variation exciting "torsional modes of synchronous generators" [ERCOT-LLWG].

The rules forming. ERCOT has proposed a **10 MW peak-to-peak limit on LCL active-power variation within any rolling 5-second window** (0.1–55 Hz evaluation band, PMU monitoring) — still a proposal in the stakeholder process as of this writing, with the revision request planned from Q2 2026 [ERCOT-LLWG]; NPPR1234 already embeds subsynchronous-oscillation requirements for large loads (effective Dec 2025) [ERCOT-1234]; NERC's May 2026 draft guideline recommends transmission owners adopt **IEEE 1453 flicker limits** into large-load interconnection requirements [NERC-GUIDE].

The mitigation market — and where it validates Zhonhen. NVIDIA's own GB300 answer is power smoothing: capped ramps, **65 joules of capacitor storage per GPU on the power shelf**, and a "GPU burn" taper — claiming 30% peak-demand reduction [NV-GB300]. A supercapacitor-shelf market has formed around the same idea (Eaton's 420 kW XLHV module, Delta's Power Capacitance Shelf, Skeleton's GrapheneGPU shelves, Musashi/Flex CESS racks) [SUPERCAP]; xAI buffers Colossus with 168+ Tesla Megapacks and Aligned deployed a 31 MW/62 MWh BESS explicitly as an interconnection asset [BESS-BUF]. **The reading (analysis): the industry has converged on Zhonhen's architectural premise — storage coupled to the power chain at multiple depths — and is assembling it piecewise from retrofit parts. Zhonhen ships it as one designed system: capacitor and battery units inside the 1 MW sidecar (150 A/ μ s load response), BBU shelves in the rack chain, battery compartments in the MVR container, campus BESS in front (with CATL, Section 4.1) — all on one DC bus [DECK].** Off-grid, the argument is stronger still: an islanded gas-fired campus has no grid inertia, so the 0.2–3 Hz swing hits the engines directly — the buffered DC backbone is generator protection (Section 5.2, Crusoe).

2.3 What this means commercially

Ride-through and flicker are the two criteria where the buying decision turns on **architecture, not brand** — a spec sheet a challenger can win. The sequencing follows: sell in ERCOT first (where the rule is live and the targets in Section 5 are building), make compliance the acceptance test, and carry the FERC timeline as the reason the rest of the country follows by 2027.

3. The field — where Zhonhen sits among SST and MV-solution providers

The frame that organizes everything: there are two competing paths from medium voltage to 800VDC, and NVIDIA has now put dates on both. Path one is the **transformer-rectifier unit (TRU)** — a line-frequency transformer plus centralized rectification, the architecture class Zhonhen productized as Panama/MVR in 2019. Path two is the **solid-state transformer (SST)** — power electronics replacing the transformer entirely. NVIDIA's August 2026 blog lays out the productization ladder it is driving with Google and Microsoft through OCP: side-mounted **power racks in 2H 2026**, centralized **row power centers in 2027**, and data-hall **DC power blocks — the step that converts MV AC directly to 800VDC, including via SST — as a 2029 target** [NV-0826, SR-0826]. An LVDC SST specification reached only v0.3 in July 2026 [NV-0826].

The strategic reading (analysis): the Western ecosystem's own roadmap reaches Zhonhen's core capability *last*. One-stage MV-to-DC conversion is the 2029 rung of NVIDIA's ladder — and it is the rung Zhonhen has been shipping commercially since the Alibaba co-development, with a claimed 15-year, 10GW+ delivered field record behind it [DECK]. That defines the window: **2026–2029, before Western SSTs productize, is when a field-proven TRU vendor is selling something nobody in the West can ship at scale.** Schneider Electric's own published architecture note says exactly this out loud: at scale, MV converts directly to 800VDC "either via a solid-state transformer (SST) or a transformer rectifier unit (TRU)" — and Schneider has announced no SST product [SE-800V].

The competitive map (as of 2026-08-22)

VENDOR	MV→DC APPROACH	STATUS	READ VS ZHONHEN
Delta	SST demo (MV AC → 800VDC, 98.5% claimed) + shipping 800V racks/shelves	SST at demo stage (OCP, Oct 2025) [DELTA]	The deepest components rival; its SST is a demo, its strength is the rack-side chain and the NVIDIA relationship
Eaton	Bought Resilient Power Systems (Austin, MV SST), Aug 2025; 800VDC reference architecture with supercaps	SST via acquisition, pre-production [EATON]	The US incumbent buying its way to MV-DC; no field record at Zhonhen's scale
Siemens + Reinhausen	Jointly developing SST, ≤36kV AC → 800VDC, ms-scale load response	Development; announced 14 Aug 2026 , no production date [SIEM]	A week-old announcement — confirms the race, and confirms nothing ships soon
Heron Power	"Heron Link" modular MW-scale SST with integrated fast-discharge Li-ion (~30s ride-through, aims to eliminate the UPS) ; \$140M Series B (Feb 2026), factory targeted 2027; on NVIDIA's roster	Pre-production [HERON]	The closest conceptual competitor — its storage-integrated MV-to-DC pitch is Zhonhen's architecture argument, restated by an ex-Tesla team with no delivered fleet
Schneider Electric	Published position: SST or TRU at scale; no SST product announced ; prefab power centers at volume	Architecture papers + prefab capacity [SE-800V]	Not a product competitor in MV-DC today — a potential channel (Section 4)
Vertiv	800VDC portfolio announced for 2H 2026; MGX reference architecture; trade reporting names Hopewind as an 800V HVDC subcontractor	Portfolio ramping [VERT, CNESA]	The Western-brand-wraps-Chinese-power-electronics model already operating in this exact market
Megmeet	800V HVDC power racks/sidecars (570kW rack at GTC 2025); on NVIDIA's components roster; SST described as roadmap in Chinese coverage	Shipping-track rack power [MEG]	The Chinese peer that got onto NVIDIA's list — proof a mainland vendor can; it plays rack-side, not MV-side
ABB / Hitachi Energy / GE Vernova / Mitsubishi	NVIDIA "data center power systems" collaborators; grid-to-rack architecture work	Collaborations, no shipping MV-DC product named [NV-1025]	ABB is also Zhonhen's eight-year strategic partner (Dec 2025 signing covering 800V for AI computing) [ZH-DOSSIER]
Kehua Data / Huawei Digital Power	China HVDC rivals; Kehua's 800V order momentum reported narrowing the gap domestically	Shipping (China); no US path	The home-market fight the 800V transition re-opened [ZH-DOSSIER]

Where Zhonhen genuinely leads — and where it is exposed

Leads (sourced): the only vendor cohort with a delivered, decade-plus MV/HVDC-for-data-centers fleet (10GW+ claimed across BAT, carriers and state operators [DECK]); one-stage MV-to-DC shipping today at 98.5% full-link efficiency, 720kW–3.6MW per system, with prefabricated MVR/MV/IT containers (IP55, C5-M, seismic-tested — export-site specs) [DECK]; a full chain from MV container down to rack silicon (1MW sidecar at 150A/μs load response, 100kW shelf, 18kW PSU, CRPS) [DECK]; lead authorship of China's national DC-power standards and the MIIT single-champion designation [ZH-DOSSIER]; TÜV Rheinland CE for AIDC power landed January 2026 [ZH-DOSSIER].

Exposed (equally sourced — carry these honestly): Zhonhen appears on **neither** NVIDIA 800VDC partner roster (May 2025: Delta, Flex, Lead Wealth, LiteOn, Megmeet in components; Oct 2025 adds ABB, Eaton, GE Vernova, Heron, Hitachi Energy, Mitsubishi, Schneider, Siemens, Vertiv in systems) [NV-0525, NV-1025] — and the company itself denied signed NVIDIA/Meta/Google agreements in August 2025 [ZH-DOSSIER]. No US entity, UL listing, or US reference site exists [ZH-DOSSIER]. **Vocabulary note (amended 2026-08-22):** "Panama Architecture" now appears by name in NVIDIA's August 2026 white paper (p22) as one of three canonical TRU implementations [NV-0826v] — the earlier warning that it was Chinese-industry shorthand absent from NVIDIA documents no longer holds. Still lead with "TRU" as the device-class term in US conversations (it is the category NVIDIA and Schneider organize by), and remember the paper names an *architecture*, not any company. The deck's strongest card — the "simpler and highly familiar design" quote — is **verified**: the white paper is in the repo ([repository-information/industry-guidance/](#)), so it can now be used with customers, cited as NVIDIA's August 2026 white paper and paraphrased (NVIDIA's notice bars verbatim reproduction).

Positioning statement (analysis, the line to carry into the call): *Zhonhen is the incumbent of an architecture the West has only just standardized around — the task is not to win the SST future, it is to monetize the 2026–2029 gap in which nothing Western ships one-stage MV-to-DC at scale, using ride-through and flicker performance (Section 2) as the spec that makes the field record legible to American buyers.*

4. The two relationships — how to spend them

4.1 CATL (public) — capital, cells, and a boundary to police

The facts. Definitive agreements signed 14 August 2026: CATL invests ~RMB 4.1B (RMB 3.512B cash plus ~49% of Times Tianyuan valued ~RMB 588M) for **49% of Zhonhen's controlling holding company**; post-deal the holdco is Zhu Guoding 35.7% / Bao Xiaoru 15.3% / CATL 49%. CATL's stated rationale: "computing power–electricity synergy" — AIDC power, EV charging/swap, new power systems and energy trading; Caixin calls it CATL's third major AI-infrastructure investment in five months [CATL-DEF, CAIXIN]. CATL's TENER Stack (9MWh, the largest single-container BESS launched, May 2025) is marketed explicitly at AI-driven data centers [TENER]. The deck's slide 7 already draws the combined offer: one source-to-rack chain — renewables/BESS/gas turbine → MV switchgear → HVDC/MVR/SST → sidecar/BBU → rack — under one EMS [DECK].

Play 1 — sell the combined power-quality machine (the technical synergy, analysis). Section 2's two buying criteria — ride-through and flicker — are both solved by the same physical thing: **storage coupled to a DC bus at multiple depths**. Zhonhen's architecture already hangs batteries directly on the DC busbar (zero-transfer ride-through), puts capacitor and battery units inside the 1MW sidecar, and offers battery compartments in the MVR container [DECK, ZH-DOSSIER]; CATL supplies the cells at every depth plus campus-scale TENER blocks in front. Heron Power's entire \$140M-funded pitch — MV-to-DC conversion with integrated fast-discharge storage that eliminates the UPS [HERON] — is this exact combination, promised for 2027+. CATL+Zhonhen can deliver it field-proven now. That is the sentence to say to Jacky: *the CATL deal turns the power-quality problem into our product demo.*

Play 2 — balance sheet and staying power. A RMB 2.1B-revenue vendor asking US buyers to bet a campus schedule on it has a credibility problem no spec sheet fixes. A vendor 49%-backed at the holdco by the world's largest battery maker can carry pilot pricing, extended warranty backstops, and multi-year service commitments that make the early-adopter deal underwritable [analysis].

The boundary (know it cold, volunteer none of it). In the US, CATL is on the DoD 1260H list (Jan 2025), NDAA-restricted for defense procurement, and CATL content strips ITC eligibility from new storage builds under the OBBBA FEOC thresholds (55% non-FEOC in 2026 → 75% by 2030) [FEOC, NDAA-1260H]. Three handling rules follow: **(a)** data-center power equipment itself claims no tax credit — DC-side ride-through storage sold as part of the power chain

is a commercial purchase, not an ITC asset; the FEOC conversation only opens where a customer's project finance wants credits on attached storage [ZH-BRIEF]; **(b)** keep every product **battery-agnostic** — the MVR container already specs lead-acid or lithium [DECK], so credit-sensitive or security-sensitive buyers can take the architecture with US or allied cells; **(c)** in US accounts, lead with the architecture and the field record — CATL is capital, China-market synergy, and cell supply where it is welcome, not the badge on the door. Where the CATL name *does* open doors is outside the US (SEA, and grid-side Gulf work) — reference sites there feed US credibility without triggering the US policy stack [ZH-BRIEF].

4.2 Schneider Electric (confidential — first-hand from Zhu Yikun) — the channel that solves all three fatal gaps

Handling first. The relationship's existence is Jacky's private disclosure. Discuss it freely *with him*; never name it to anyone else; outside Zhonhen conversations, frame the topic as a hypothetical about "partnering with a US electrical incumbent."

Why Schneider specifically is the right door (all public, safe to reason from). Schneider is tied with Vertiv for #1 in data-center physical infrastructure globally (Dell'Oro) [DO-25]; it signed a \$1.9B supply-capacity agreement with Switch for prefabricated power modules (Nov 2025) [SE-SW]; it co-publishes validated AI-factory blueprints with NVIDIA and runs the ETAP grid-to-chip digital twin [SE-NV]; it is spending \$700M through 2027 on US manufacturing, including new MV capacity in Texas and Tennessee [SE-MFG]; and its Red Oak, TX integration plant is built to turn out 240–400 prefab power centers a year [SE-PREFAB]. Against all that, its published 800VDC position concedes the MV-to-DC conversion will be done by "SST or TRU" — and it has no SST product and no announced TRU line [SE-800V]. **Schneider has published an architectural need for precisely the box Zhonhen ships** [analysis].

The plays, ranked by value:

1. **OEM / white-label the TRU inside Schneider's prefab power centers.** Zhonhen MVR blocks integrated at Red Oak, wearing Schneider's brand, UL pathway, and service network. This solves Zhonhen's three fatal US gaps in one structure — brand/security optics, certification, and the 3AM-service question — while giving Schneider a shipping MV-to-DC product two to three years before any Western SST reaches production [analysis]. The model is not hypothetical: trade reporting already describes Vertiv shipping 800V HVDC built on a Chinese subcontractor (Hopewind) [CNESA]. Expect Schneider to demand air-gapped, wired-management controls (which also keeps the product outside the FCC connected-inverter rule's connectivity prong [ZH-BRIEF]) and, plausibly, US final assembly — which its Texas/Tennessee expansion could host [SE-MFG].
2. **Component supply into Schneider's 800VDC reference designs** — rectifier modules, DC/DC stages, shelves. Lower value, faster to start, and a proof-of-quality on-ramp to play 1.
3. **Certification and test partnership** — using Schneider's US lab muscle to compress the UL roadmap that currently gates every direct sale [ZH-BRIEF].

What not to expect (analysis): Schneider will not lend its customer relationships to a Zhonhen-branded sale, and any structure will be built so Schneider owns the customer. That is the acceptable price: in 2026 the alternative to being inside a US incumbent's box is being outside every US door. The question worth asking Jacky — it shows strategic judgment rather than curiosity — is *which layer* the Schneider conversation is at: components, integrated blocks, or certification. The answer determines the entire shape of the US go-to-market.

5. The neocloud thesis — verified, refined, and named

5.1 The verification verdict

Jacky's claim — *neoclouds are the target because they willingly procure power solutions, prioritizing speed and technical capability* — is **true, with one refinement that changes who you call**. "Neocloud" covers two different buyer types:

- **Owner-operator self-builders** — they own land, substations, and buildings, and buy power equipment directly with documented deals: **IREN** (owns its ERCOT substations; its January 2025 business update itemizes direct procurement of 345/138kV transformers, circuit breakers, UPS, gensets, and CDUs — with **no named electrical vendors on record**, i.e., open supplier slots) [IREN-BU]; **Crusoe** (as developer: 29 GE Vernova aeroderivative turbines ≈1GW booked 2024–25, chosen explicitly because combined-cycle lead times ran ~7 years — speed-over-incumbency procurement in its purest form) [GEV-CRU]; **xAI** (~35 turbines, 300+ Tesla Megapacks, its own substation refurbishments, a turbine JV with Solaris — and the 122-day Colossus bring-up as the sector's speed benchmark) [XAI]. The claim is fully supported here.
- **Tenant neoclouds** — **CoreWeave, Lambda, Nebius** lease from landlords who buy the MV equipment (CoreWeave's Core Scientific acquisition **failed** — shareholders rejected it 30 Oct 2025, so CoreWeave remains overwhelmingly a tenant; Lancaster, PA is its first self-build) [CORZ-VOTE]. They are *not* MV buyers — **but** NVIDIA publicly names CoreWeave, Lambda and Nebius among the "industry pioneers designing for 800-volt data centers" [NV-OCF], which makes them buyers of the *rack-side* chain (power shelves, sidecars, BBUs) that travels with the GPU deployment.
- **The adjacent landlord cohort** — Applied Digital, TeraWulf, Core Scientific — behaves exactly like the self-builder neoclouds (direct procurement, schedule pressure from anchor leases) and should be worked with the same motion [ZH-COVERAGE].

5.2 The targets, ranked

Target 1 — IREN. *Why:* the most vertically integrated buyer in the cohort — it owns everything from the 345kV substation to the liquid-cooled rack, and it buys all of it directly with no entrenched electrical vendor on public record [IREN-BU]. Its power is **ERCOT-concentrated** — Childress (750MW secured, where Horizon 1 was announced in March 2025 and energized within ~9 months) and Sweetwater (2.75GW across two phases, energization from April 2026) — which makes Section 2's ride-through requirement a *first-order procurement criterion*, not background noise [IREN-BU, ZH-COVERAGE]. And it is a proof-driven technical buyer: Horizon 1 was delivered to Microsoft and accepted on 13 August 2026 with NVIDIA Exemplar Cloud status on GB300 NVL72 [ZH-COVERAGE]. *How:* pitch a **pilot MVR container (or 240/400V HVDC block) on one Sweetwater building** — the buildout through 2028 is the procurement window, and the \$9.7B Microsoft contract makes schedule the thing money can't fix. Lead with delivery-in-weeks against the 30–40-month domestic transformer queue, ERCOT ride-through by architecture, air-gapped wired-management spec, embedded engineers, extended warranty, pilot pricing.

Target 2 — Crusoe. *Why:* the self-described "energy-first AI factory company" builds **off-grid and behind-the-meter** — Abilene's 1.2GW runs on its own generation; a ~750MW Bergen reciprocating-engine order and a 900MW Microsoft campus (with MV-connected batteries) follow; Childress with Lancium starts construction Q3 2026 — ERCOT again [ZH-COVERAGE, GEV-CRU]. *The analytical hook:* the flicker problem is **worse off-grid than on-grid** — an islanded microgrid has no system inertia, so synchronized GPU load swings hit the engines directly, forcing derating, wasted spinning reserve, and mechanical wear. A DC backbone with storage absorbed at every depth (busbar batteries, sidecar CBUs, container BESS) lets the gensets see a smoothed load. *How:* sell the **DC power backbone for gas-fired AI factories** — MV container + MVR + buffered DC bus — pitched as generator protection and fuel efficiency, not as a

UPS replacement. Crusoe's procurement culture (aeroderivatives over combined-cycle purely for schedule) is exactly the buyer that trades incumbency for capability [GEV-CRU].

Target 3 (same motion, second wave) — the brownfield landlords: Core Scientific and TeraWulf. Converted mining sites carry existing MV infrastructure and brutal retrofit schedules against anchor leases (Core Scientific: ~530MW AMD anchor, Denton and Hunt County TX with approved ERCOT interconnection; TeraWulf: Fluidstack/Google-backstopped and Anthropic leases) [ZH-COVERAGE]. The pitch is the retrofit: one-stage MV-to-DC blocks and the no-shutdown upgrade path, into halls whose power chain is being rebuilt anyway.

Deprioritize — xAI. The best technical fit on paper (largest behind-the-meter storage buffer fleet, fastest builder, biggest flicker problem) is the worst political fit in practice: xAI is now inside SpaceX — a defense contractor mid-IPO — and a Chinese MV vendor inside that supply chain is a security-review hard stop, not a negotiation [XAI, ZH-BRIEF]. Say this honestly in the call; it demonstrates judgment.

The rack-side wedge — CoreWeave, Lambda, Nebius. Not MV buyers, but 800V design pioneers [NV-OCP]: the sidecar/shelf/PSU chain (1MW sidecar at 150A/μs, 100kW shelf, 18kW PSU [DECK]) can travel with their GPU deployments into landlord halls — a second product motion on Megmeet's model, not a substitute for the MV sale. Lambda's first self-developed Kansas City AI factory (24MW → 100MW+, announced Oct 2025) is the one small direct MV opening in this group [LAMBDA-KC].

5.3 The common entry motion (analysis)

1. **Sell the schedule first** — factory-built MV-to-DC skids against a 30–40-month domestic transformer queue and scarce site electricians; the tariff conversation loses to the calendar conversation in 2026 procurement [ZH-BRIEF].
2. **Sell power quality as the spec** — the ERCOT ride-through requirement and GPU flicker (Section 2) are the two criteria where Zhonhen's architecture wins on physics, not on price. Make the pilot's acceptance test a ride-through and load-step demonstration.
3. **De-risk visibly** — pilot pricing, embedded engineering, extended warranty backstops (CATL balance sheet, Section 4.1), air-gapped controls as a spec commitment, and a named US service partner (the ABB relationship today; the Schneider structure if it matures).
4. **One lighthouse beats breadth** — a single accepted ERCOT pilot at IREN or Crusoe converts the China field record into an American reference; until then every deal is missionary work. Sequence accordingly.

6. The risk stack — what can kill this, and the counters

The full US regulatory treatment is in the interview brief ("The regulatory stack") — this table holds only what bears directly on the strategy above.

RISK	REALITY CHECK	COUNTER
FCC connected-inverter rule (eff. 28 Jul 2026)	Covers <i>bi-directional</i> connected inverters; unidirectional AC-to-DC rectification on a plain reading sits outside scope — but V2G chargers are bidirectional by definition and battery-hybrid UPS variants sit near the line [ZH-BRIEF, FCC]	Scope the US roadmap around it from day one: air-gapped, wired-management controls as a <i>published spec commitment</i> ; keep V2G and bidirectional variants out of the US catalogue
Tariffs	Raise landed cost; don't bar the door	The calendar argument: a tariffed, factory-built MV-to-DC skid available in weeks beats a 30–40-month domestic transformer queue; model landed cost per deal [ZH-BRIEF]

RISK	REALITY CHECK	COUNTER
NVIDIA-roster absence	Zhonhen is on neither the May 2025 nor Oct 2025 partner list, and denied signed hyperscaler agreements in Aug 2025 [NV-0525, NV-1025, ZH-DOSSIER]	Never imply membership. The architecture-validation claim survives on its own: NVIDIA's public materials confirm the TRU path; the once-gated whitepaper quote is now verified in hand (p22) [NV-0826v] — cite the paper, paraphrase the phrase, and never imply roster membership from it
NOGRR282 non-retroactivity	Loads approved pre-14 Nov 2025 are exempt — the compliance sale doesn't bind there [ERCOT-SPWG]	Target <i>new</i> buildings and post-2025 energizations (IREN Sweetwater phases, Crusoe Childress); sell exempt sites the voluntary-retrofit and flicker story instead
Security/procurement politics	No entity-list issue exists, but a Chinese MV vendor inside US AI infrastructure means security questionnaires and, for federal-adjacent buyers, hard stops [ZH-BRIEF]	Segment ruthlessly (Section 5): commercial neoclouds and landlords first, xAI/federal-adjacent never; lead with the Schneider structure where it matures
Governance optics	The founder's Dec 2025 conviction and the CATL 49% are one search away for any US counterparty [ZH-DOSSIER]	Precision, volunteered by no one: CATL holds 49% of the <i>holding company</i> , not the listco; Bao Xiaoru chairs; the conviction is never raised by you in any conversation — with Jacky or anyone else
Execution: no US service network	The 3AM question loses deals on its own [ZH-BRIEF]	Partner-led service from day one — the ABB relationship now, the Schneider structure as the destination; embedded engineers on every pilot

7. The call playbook

Five things to say that prove strategic command

1. *"The West's own roadmap reaches your core product last — NVIDIA's MV-to-DC power block is the 2029 rung of its ladder, and you've been shipping that rung since the Alibaba co-development. The strategy isn't to win the SST future; it's to monetize 2026 through 2029."*
2. *"ERCOT made ride-through law on August 1 — Large Computational Loads over 75 megawatts now must ride the sag envelope, recover to 90% within two seconds, and disturbance-counting trips are banned outright. FERC ordered national standards by December. Batteries on the DC busbar means we comply by construction — I'd make the NOGRR282 envelope the acceptance test of every pilot." (Precision points: ERCOT's record is 26 events over 100 MW; the 1.5 GW and 3 GW mass-trips were Northern Virginia and PJM.)*
3. *"The flicker problem is your second proof: Microsoft, OpenAI and NVIDIA published that training swings run tens to hundreds of megawatts at 0.2 to 3 hertz — turbine torsional territory — and the industry's answer is capacitor shelves and Megapacks bolted onto AC architectures. We ship storage at four depths on one DC bus as a designed system. And ERCOT has a 10-megawatt-per-5-seconds variation limit in the pipeline."*
4. *"On neoclouds you're right, and I'd sharpen it: the self-builders are the MV buyers. IREN owns its ERCOT substations and lists no electrical vendor publicly — open slots, 2.75 gigawatts building at Sweetwater. Crusoe builds off-grid, where GPU swings hit the gensets directly — our buffered DC backbone is generator protection there. CoreWeave, Lambda and Nebius lease, so they get the sidecar-and-shelf motion, not the MV motion. And xAI is inside SpaceX now — that door is closed to us, and saying so saves us a year."*
5. *"The CATL deal isn't just capital — it completes the power-quality machine your slide 7 already draws. But in America it stays behind the architecture: battery-agnostic products, FEOC-clean options, CATL as balance sheet and cells where they're welcome — never the badge on the door."*

The Schneider conversation (with Jacky only)

Since he raised it, engage it at the strategy level: *"Schneider's own 800VDC paper says the MV conversion happens by SST or TRU — and they have neither product. Their Red Oak plant integrates prefab power centers at hundreds a year, and they're spending \$700 million on US factories. A Zhonhen TRU inside a Schneider-branded power center solves our brand, certification and service problems in one structure — Vertiv is already reported to ship 800V HVDC built on a Chinese subcontractor. The question that decides our US shape: is the Schneider conversation at the component layer, the integrated-block layer, or certification?"* Outside conversations with Zhonhen, the relationship does not exist.

Questions worth asking him

1. Which layer is the Schneider engagement at — components, integrated blocks, or certification — and who inside Zhonhen owns it?
2. What is the real UL/US-certification roadmap and date? (It gates every direct deal; TÜV CE landed January 2026 — what's next?)
3. Can the company underwrite pilot economics — extended warranty backstops, embedded engineers — and does the CATL relationship change what we can offer?
4. Is anyone building the ERCOT compliance artifact — a documented NOGRR282 ride-through test on an MVR container? If not, can I make that my first project?
5. For the US, do we commit to the "TRU / MV rectifier" vocabulary and an air-gapped controls spec as published positions?

Do not say

- Anything about the founder's conviction — to Jacky or anyone.
- "The 108-kilowatt shelf" — say **"the 100-kilowatt power shelf"** (slide 6 vs slide 15 discrepancy).
- Any claim that Zhonhen is an NVIDIA partner or on the 800VDC roster — it is not, and the company itself denied signed agreements in August 2025.
- Implying roster membership from the whitepaper's Panama mention — the quote itself is now verified in hand (p22 [NV-0826v] and usable when cited and paraphrased, but it endorses an architecture class, never a company.
- "Panama" to an American buyer — it's "MV rectifier" or "TRU" in every US-facing sentence.
- The Schneider relationship, anywhere outside Zhonhen.

8. Sources

Internal documents (the foundation):

- **[ZH-DOSSIER]** — `zhonhen.profile.json`, profileVersion 1 (2026-08-21), 41 public sources — company shape, financials, HVDC share (Kezhi 2025), Enervell/SuperX, governance record
- **[DECK]** — `zhonhen-deck-summary.md` — the company's confidential AIDC introduction deck v1.35 (2026-08-19); all deck-sourced specs and claims are company-stated and confidential
- **[ZH-BRIEF]** — `zhonhen-interview-brief.md` — the US regulatory stack, objection handling, tariff/schedule math
- **[ZH-COVERAGE]** — Profiler dossiers: `iren`, `crusoe`, `coreweave`, `core-scientific`, `terawulf`, `xai` profile JSONs (2026-08-21/22 batch), each with its own sources
- **[FLAG]** — verification flag: the NVIDIA "Industry Alignment" whitepaper PDF is access-gated; the "simpler and highly familiar design" quote and the "Panama Architecture" naming could not be publicly verified as of 2026-08-22

Ride-through and flicker (Section 2):

- **[NERC-0125]** — NERC Incident Review, "Considering Simultaneous Voltage-Sensitive Load Reductions," Jan 2025 — https://www.nerc.com/globalassets/our-work/reports/event-reports/incident_review_large_load_loss.pdf
- **[NERC-0126]** — NERC Incident Review, "Voltage-Sensitive Crypto Load Reductions," Jan 2026 (the 26 ERCOT events; trip-chain physics) — https://www.nerc.com/globalassets/our-work/reports/event-reports/incident_review_considering_voltage-sensitive_crypto_load_reductions.pdf
- **[PJM-0726]** — POWER Magazine, "3 GW Off PJM, Again," 2026-08-10 — <https://www.powermag.com/3-gw-off-pjm-again-what-nercs-2024-investigation-already-told-us/>
- **[ERCOT-1308]** — ERCOT NOGRR282 + NPPR1308 (LCL ride-through; PUCT-approved 2026-07-09, effective 2026-08-01) — <https://www.ercot.com/mktrules/issues/NOGRR282> · <https://www.ercot.com/mktrules/issues/NPPR1308>
- **[ERCOT-SPWG]** — ERCOT (Gravois), "Large Electronic Load Ride-Through Requirements," SPWG deck, 2026-02-24 (curve values are from this pre-approval draft; confirm final language before external citation) — https://www.ercot.com/files/docs/2026/03/02/04_LEL-RT-Requirements_SPWG_Feb2026.pdf
- **[ERCOT-BOD]** — ERCOT Board Item 6.2 (Woodfin memo), 2025-12-08 — <https://www.ercot.com/files/docs/2025/12/01/6.2-NOGRR282-Large-Electronic-Load-Ride-Through-Requirements-and-NPPR1308.pdf>
- **[ERCOT-1234]** — ERCOT NPPR1234, large-load interconnection/modeling incl. subsynchronous-oscillation requirements (fully effective 2025-12-15) — <https://www.ercot.com/mktrules/issues/NPPR1234>
- **[SB6]** — Texas SB6 (signed 2025-06-21): McGuireWoods alert, Jul 2025 — <https://www.mcguirewoods.com/client-resources/alerts/2025/7/texas-senate-bill-6-significantly-expands-regulatory-oversight-over-large-loads-in-ercot/>
- **[NERC-ALERT]** — NERC Level 3 Essential Action alert, 2026-05-04 (via MRO/Utility Dive) — <https://www.mro.net/nerc-issues-critical-reliability-alert-as-emerging-large-loads-challenge-the-grid/>
- **[FERC-RD26]** — FERC Docket RD26-7-000, order directing mandatory NERC computational-load standards, 2026-07-16 (via POWER) — <https://www.powermag.com/ferc-orders-mandatory-nerc-reliability-standards-for-data-center-and-other-computational-loads/>
- **[MISO-LL]** — MISO Large Load Interconnection Reliability Requirements initiative — <https://www.misoenergy.org/engage/MISO-Dashboard/large-load-interconnection-reliability-requirements/>
- **[NERC-LLTF]** — NERC LLTF white paper, "Characteristics and Risks of Emerging Large Loads," Jul 2025 (the 1.9 p.u./s recording) — <https://www.nerc.com/globalassets/who-we-are/standing-committees/rstc/whitepaper-characteristics-and-risks-of-emerging-large-loads.pdf>
- **[NERC-GUIDE]** — NERC draft Reliability Guideline, "Risk Mitigation for Emerging Large Loads," May 2026 (IEEE 1453 recommendation) — https://www.nerc.com/globalassets/who-we-are/standing-committees/rstc/draft_reliabilityguideline_riskmitigationforemerginglargeloads.pdf
- **[ARXIV-PS]** — Microsoft/OpenAI/NVIDIA, "Power Stabilization for AI Training Datacenters," arXiv:2508.14318, Aug 2025 — <https://arxiv.org/abs/2508.14318>
- **[LLAMA]** — Meta, "The Llama 3 Herd of Models," arXiv:2407.21783, Jul 2024 (tens-of-MW fluctuation passage) — <https://arxiv.org/abs/2407.21783>
- **[ERCOT-LLWG]** — ERCOT LLWG, "LCL Power Variation Limit and Monitoring Requirement," 2026-04-24 (the proposed 10 MW / 5-s limit — a proposal, not adopted, as of 2026-08-22) — https://www.ercot.com/files/docs/2026/04/24/ERCOT-Large-Computational-Load-LCL-Power-Variation-Limit-and-Monitoring-Requirement-Need-and-Discussion_fin.pdf

- **[NV-GB300]** — NVIDIA Technical Blog, GB300 NVL72 power smoothing (65 J/GPU, 30% peak reduction), 2025-07-28 — <https://developer.nvidia.com/blog/how-new-gb300-nvl72-features-provide-steady-power-for-ai/>
- **[SUPERCAP]** — IEEE Spectrum supercapacitor roundup, 2025-05-06 (Eaton XLHV, Siemens Energy E-statcom, Delta PCS) — <https://spectrum.ieee.org/supercapacitor-2671883490> ; Skeleton GrapheneGPU — <https://www.skeletontech.com/news/skeletons-new-product-cuts-ai-data-center-energy-use-by-44-and-boosts-computing-power-by-40> ; Flex/Musashi CESS — <https://investors.flex.com/news/news-details/2024/Flex-and-Musashi-Energy-Solutions-Partner-to-Mitigate-Utility-Power-Challenges-in-AI-Data-Centers/default.aspx>
- **[BESS-BUF]** — xAI Megapacks: DCD, 2025 — <https://www.datacenterdynamics.com/en/news/xai-deploys-168-tesla-megapacks-to-power-its-colossus-supercomputer-in-memphis/> ; Aligned/Calibrant 31 MW BESS, 2025-10-22 — <https://www.powermag.com/aligned-calibrant-deploy-battery-storage-to-support-data-centers/>
- **[FCC]** — FCC power-inverter Covered List determination, effective 2026-07-28 — <https://www.fcc.gov/sites/default/files/power-inverter-fcc-determination.pdf>

The 800VDC ecosystem and competitors (Section 3):

- **[NV-0525]** — NVIDIA Technical Blog, 800VDC announcement + first partner roster, 2025-05-20 — <https://developer.nvidia.com/blog/nvidia-800-v-hvdc-architecture-will-power-the-next-generation-of-ai-factories/>
- **[NV-1025]** — NVIDIA Technical Blog, "Building the 800 VDC Ecosystem" (expanded roster incl. Heron Power), 2025-10-13 — <https://developer.nvidia.com/blog/building-the-800-vdc-ecosystem-for-efficient-scalable-ai-factories/>
- **[NV-0826]** — NVIDIA Blog, "Why Scaling AI Compute Performance Requires a New Power Architecture" (product ladder: racks 2H26, row centers 2027, DC power blocks/SST 2029; LVDC SST spec v0.3), 2026-08-11 — <https://blogs.nvidia.com/blog/800-vdc-power-architecture-ai-factory/>
- **[NV-OCP]** — NVIDIA corporate blog, OCP 2025 ("CoreWeave, Lambda, Nebius, Oracle Cloud Infrastructure and Together AI... designing for 800-volt data centers"), 2025-10-13 — <https://blogs.nvidia.com/blog/gigawatt-ai-factories-ocp-vera-rubin/>
- **[SR-0826]** — StorageReview coverage of the NVIDIA Industry Alignment whitepaper (Options A/B/C, SST 2029 target), 2026-08-15 — <https://www.storagereview.com/news/nvidia-moves-800-vdc-power-architecture-from-concept-to-production-just-dont-turn-off-ac-power-yet>
- **[DELTA]** — Delta OCP 2025 SST/HVDC demonstrations — <https://www.deltapowersolutions.com/en/mcis/news-2025-delta-demonstrate-hvdc-power-cooling-networking-to-drive-ai-data-center-evolution-at-ocp-global-summit.php>
- **[EATON]** — Eaton 800VDC reference architecture + Resilient Power Systems (MV SST) acquisition, 2025 — <https://www.eaton.com/us/en-us/company/news-insights/news-releases/2025/eaton-unveils-next-generation-architecture.html>
- **[SIEM]** — Siemens + Reinhausen SST partnership (≤ 36 kV $\rightarrow$ 800VDC), announced 2026-08-14 — <https://www.engineering.com/siemens-and-reinhausen-partner-on-solid-state-transformer/>
- **[HERON]** — Heron Power: \$140M Series B, 2026-02-18 (TechCrunch) — <https://techcrunch.com/2026/02/18/heron-power-raises-140m-to-ramp-production-of-grid-altering-tech/> ; product blueprint — <https://heronpower.com/blueprint-for-800-vdc>
- **[VERT]** — Vertiv 800VDC portfolio for 2H 2026, 2025-10 — <https://www.prnewswire.com/news-releases/from-vision-to-readiness-vertiv-collaborates-with-nvidia-to-advance-800-vdc-platform-designs-to-power-the-next-generation-of-ai-factories-302582190.html>
- **[MEG]** — Megmeet 800V HVDC racks (570 kW, GTC 2025) — <https://eu.36kr.com/en/p/3234155775573637>
- **[CNESA]** — China Energy Storage Alliance, "NVIDIA's 800V Architecture... 31 Core Industry Chain Companies," 2025-11-25 (the Vertiv/Hopewind subcontractor reporting) — <https://en.cnesa.org/latest-news/2025/11/25/nvidias->

800v-architecture-reshapes-ai-data-centers-31-core-industry-chain-companies-unveiled

Schneider Electric (Section 4.2, public record):

- **[DO-25]** — Dell'Oro: Schneider/Vertiv virtually tied for #1 in DC physical infrastructure, 2025 — <https://www.cablinginstall.com/data-center/article/14291375/schneider-electric-vertiv-square-off-on-data-center-physical-infrastructure-analyst>
- **[SE-800V]** — Schneider Electric 800VDC position ("either via a solid-state transformer (SST) or a transformer rectifier unit (TRU)") — <https://www.se.com/www/en/insights/ai-and-technology/artificial-intelligence/vdc-powering-the-future-of-ai-data-centers/>
- **[SE-SW]** — Schneider–Switch \$1.9B supply-capacity agreement, 2025-11-19 — <https://www.prnewswire.com/news-releases/schneider-electric-and-switch-expand-partnership-with-1-9-billion-supply-capacity-agreement-to-power-ai-factories-302620152.html>
- **[SE-NV]** — Schneider–NVIDIA validated gigawatt AI-factory blueprints (GTC 2026-03-16) — <https://www.globenewswire.com/news-release/2026/03/16/3256775/0/en/Schneider-Electric-teams-with-NVIDIA-to-develop-validated-blueprints-to-design-simulate-build-operate-and-maintain-gigawatt-scale-AI-Factories.html>
- **[SE-MFG]** — Schneider \$700M US manufacturing investment through 2027 — <https://www.manufacturingdive.com/news/schneider-electric-700-million-investment-us-data-center-demand-ai/743630/>
- **[SE-PREFAB]** — Schneider Red Oak, TX prefab integration facility (240–400 power centers/yr) — <https://www.datacenterdynamics.com/en/product-news/schneider-electric-opens-data-center-integration-facility-in-red-oak-texas/>

CATL (Section 4.1):

- **[CATL-DEF]** — Definitive CATL–Zhonhen holdco investment (49%, ~RMB 4.1B), signed 2026-08-14 — <https://www.energytrend.com/news/20260817-51953.html>
- **[CAIXIN]** — Caixin Global, "CATL Buys Into AI Data Center Supply Chain," 2026-08-18 — <https://www.caixinglobal.com/2026-08-18/catl-buys-into-ai-data-center-supply-chain-as-it-seeks-next-growth-market-102475497.html>
- **[TENER]** — CATL TENER Stack 9 MWh launch (marketed at AI data centers), 2025-05 — <https://www.catl.com/en/news/6410.html>
- **[FEOC]** — OBBBA FEOC material-assistance thresholds (Morgan Lewis, Mar 2026) — <https://www.morganlewis.com/pubs/2026/03/how-feoc-rules-are-reshaping-energy-storage-tax-credit-eligibility>
- **[NDAA-1260H]** — CATL added to DoD 1260H list, Jan 2025 — https://www.upi.com/Top_News/US/2025/01/07/Tencent-CATL-Chinese-military-companies/6101736233952/

Neocloud procurement (Section 5):

- **[IREN-BU]** — IREN Business Update (direct procurement of substation transformers, breakers, UPS, gensets, CDUs), 2025-01-21 — <https://www.globenewswire.com/news-release/2025/01/21/3012652/0/en/IREN-Business-Update.html>
- **[GEV-CRU]** — GE Vernova–Crusoe 29-unit aeroderivative turbine order (≈1 GW), 2025-06 — <https://www.gevernova.com/news/press-releases/ge-vernova-crusoe-announce-major-29-unit-aeroderivative-gas-turbine-deliver-ai-data-centers>
- **[XAI]** — xAI Colossus procurement (turbines, Megapacks, Solaris JV; 122-day bring-up): SemiAnalysis — <https://newsletter.semianalysis.com/p/xais-colossus-2-first-gigawatt-datacenter> ; DCD Megapacks (see [BESS-BUF])

- **[CORZ-VOTE]** — Core Scientific shareholders reject CoreWeave's \$9B acquisition; merger terminated, 2025-10-30 (CNBC) — <https://www.cnbc.com/2025/10/30/core-scientific-shareholders-reject-9-billion-coreweave-offer-deal-terminated.html>
- **[LAMBDA-KC]** — Lambda's first self-developed AI factory (Kansas City, 24 MW → 100 MW+), 2025-10-28 — <https://www.businesswire.com/news/home/20251028473000/en/Lambda-Doubles-Down-on-Midwest-Expansion-To-Build-AI-Factory-in-Kansas-City-MO>

Zhonhen Electric — US AIDC Strategy Report · Internal & confidential, strategy preparation · Rendered from zhonhen-strategy-report.md, which remains the source of truth · Developed by: LightAISolutions